
Can Power Draw Constrain Covert Compute?

Tom Kimpson¹ Mauricio Baker² Emlyn Graham³ Anjay Friedman² Maximus Rafla² Coby Joseph²

Abstract

Managing the risks of frontier AI will require international agreements.

Agreements require verification: an outside party must be able to check what a compute operator actually ran, for example constraining the total number of floating-point operations (FLOP).

Verification must be robust to an adversarial operator who hides computation, yet the hardware mechanisms proposed to catch one typically read silicon the operator controls.

Analogue side-channel measurements have been identified as a path forwards: sensors for power, electromagnetic, acoustic, or thermal emissions sit off the chip, beyond the operator’s reach.

We quantify how tightly an off-chip power meter can constrain compute.

We derive a closed form for this tightness: the floor β , the largest covert computation a power trace cannot rule out, in units of one machine’s declared capacity.

We measure every quantity entering the closed form on NVIDIA A100 GPUs.

We find the floor is large: we construct explicit cheating strategies that attain $\beta = 0.4$ behind a trace the meter accepts, and the closed form bounds the worst case at $\beta = 1.16$ when the declared number format can be checked, rising to $\beta = 1.95$ against a bare meter.

Additional verifier capabilities each buy a stated reduction — re-executing the declared work at an observed operating point brings β to 0.34. However no capability reaches the covert side, which closes only by assumption on the adversary.

A power-based verification regime can then be designed or rejected... etc.

Dev note: abstract is a placeholder — TBC once the main text is locked in. Deliberately left to last.

1. Introduction

As frontier AI systems grow more capable (Sevilla et al., 2022; Bengio et al., 2026), governance proposals increasingly rely on *verification*: the ability of an outside party to check what a monitored operator is doing (Reuel et al., 2025; Baker et al., 2025). One proposed mechanism is a threshold on compute (Sastry et al., 2024). Above a specified number of floating-point operations (FLOP), a training run triggers registration, reporting, or audit obligations (Shavit, 2023; Baker et al., 2025; Heim & Koessler, 2024). For example, U.S. Executive Order 14110 temporarily required reports on models trained using more than 10^{26} integer or floating-point operations (The White House, 2023; 2025).

Compute verification policies need to be robust to adversarial operators who do not behave honestly, but instead try to subvert the threshold by hiding computation. Verification methods that rely on passive telemetry such as utilisation counters, on-chip clocks or memory traffic read silicon the operator controls. Consequently they rely on tamper-resistant read paths and attested monitoring firmware that do not yet exist (O’Gara et al., 2025; Ansari, 2026; Rahman & Tajdari, 2026). Cryptographic proof-of-training protocols (Jia et al., 2021; Garg et al., 2023) depend on the operator’s cooperation. Active challenge–response schemes (Monfared et al., 2026) run their probes on the operator’s own untrusted hardware.

“Off-chip” analogue sensors have been suggested as an alternative verification method (Baker et al., 2025). This involves measurement hardware physically separate from the monitored chips such as a meter on the power feed, an antenna beside the rack, a microphone, or a thermal camera. Each reads a physical side channel such as power, electromagnetic, acoustic, or thermal emissions and so its measurements reach the verifier without passing through any silicon or software the operator controls. Such sensors could be retrofitted to existing data centres and physically secured against tampering by the adversary.

There are reasons for optimism: in non-adversarial settings, physical emissions are known to leak workload information.

¹School of Mathematics and Statistics, The University of Melbourne, Parkville, VIC, Australia ²RAND Corporation ³Affiliation TBD. Correspondence to: Tom Kimpson <tom.kimpson@unimelb.edu.au>.

Emissions let an observer recover model architectures and hyperparameters (Hu et al., 2020; Maia et al., 2022; Gao et al., 2023; Chaudhuri et al., 2025), and software-visible counters and timing let a co-resident process do the same (Naghijouybari et al., 2018; Wei et al., 2020). On-node telemetry and on-chip counters classify workloads (Tang et al., 2022; Xu et al., 2026; Gangwal et al., 2019), and power draw betrays anomalous computation (Clark et al., 2013) and running GPU activity (Arefin & Serwadda, 2024); a related body of work characterises power draw for GPUs and data centres (Patel et al., 2023; Choukse et al., 2025; Latif et al., 2025). However these are forward models, fitted largely in cooperative conditions, i.e. they map known workloads to their emissions. Verification poses the inverse problem (i.e. emissions to workload) against an operator actively trying to evade the sensor. It is an open question how well analogue-sensor measurements can constrain compute in that setting. This paper addresses that question for one channel: power, the channel most directly linked to computation, since every operation costs energy and total energy therefore budgets total work. We defer the study of other channels to a future work.

This paper is organised as follows. Section 2 sets up the problem mathematically. Section 3 derives a floor on the covert compute a power trace admits, in closed form, from a handful of physical bands. Section 4 empirically determines how large that floor is by measuring the bands directly on NVIDIA A100 hardware. Section 5 shows how additional information gained by the verifier beyond the power trace (e.g a precision declaration) pushes the floor lower. We close with a discussion in Section 6.

2. Setup and threat model

2.1. The forward model

Fix an observation window $t \in [0, T]$ and let $F(t)$ be the instantaneous rate of arithmetic operations, bounded by the hardware capacity $F_{\max}$. The forward model of $F(t)$ to power follows from the CMOS power law (Chandrakasan et al., 1992); chip power splits into a switching term and a static term,

$$P = \underbrace{\alpha C V^2 f}_{\text{dynamic}} + P_{\text{stat}}, \quad (1)$$

with V the supply voltage, f the clock frequency, C the total capacitance of the logic, and α the activity factor: the fraction of C that actually toggles in a given cycle ($0 \leq \alpha \leq 1$), so that αC is the capacitance switched per clock tick. A datapath that retires N_{op} nominal operations per cycle runs at operation rate $F = N_{\text{op}} f$. Substituting $f = F/N_{\text{op}}$ and collecting constants gives the forward model

$$P(t) = r(t) F(t) + P_0(t), \quad r \equiv \alpha \left(\frac{C}{N_{\text{op}}} \right) V^2, \quad (2)$$

where P_0 collects P_{stat} and the remaining non-compute overhead (idle draw, cooling, I/O). We call r the *exchange rate* [J/op]. The clock frequency has dropped out: what an operation costs does not depend on how fast the operations run, only on what each switching event costs and how many operations it delivers. Appendix A gives this derivation in full.

The exchange rate $r(t)$ is not a constant of the device but a property of the execution, and it changes whenever the execution does. Take its three factors in turn. The activity factor α depends on the operand values being multiplied (Chandrakasan et al., 1992). The capacitance switched per operation, C/N_{op} , depends on the numeric format (which sets both the datapath width and the packing N_{op}) and on where the operands reside (Horowitz, 2014). The supply voltage V tracks the voltage–frequency operating point (Palipadi & Starikovskiy, 2006). Every one of them is under the prover’s control, and none is visible to an off-chip meter. Empirically the spread is wide: component-level analyses of a 45 nm process put arithmetic and on-chip memory access at 0.1–20 pJ per operation depending on format and operand locality (Horowitz, 2014).

In practice, the meter reports only a noisy version of the power,

$$P_{\text{obs}}(t) = P(t) + \eta(t), \quad (3)$$

with $\eta(t)$ a zero-mean measurement noise the prover cannot forge. The noise is a measurement limit: zero-mean noise with a suitable correlation structure (stationary, finitely correlated) averages away over the window and shrinks the longer the verifier watches. Throughout this paper we therefore assume a perfect meter and work with the clean power $P(t)$; every floor we derive is a floor on what even a noiseless meter can establish.

We count hardware operations (i.e. what the silicon actually executes) rather than model FLOPs (i.e. the algorithmically required operations). Every executed operation costs joules, so the meter sees hardware operations. We count them as nominal operations of whichever numeric format runs: a multiply–add retires two nominal ops whether it executes at fp32 or int8, and integer multiply–adds are ops, not FLOPs.

2.2. The threat model

There are two parties. The *prover* is the monitored operator: it controls the workload, the operand values, the numeric formats, the clocks, and the idle states of the machine, and it may wish to run compute it does not declare. The *verifier* is the inspecting party: it reads P_{obs} from an off-chip meter it trusts, and it never sees r , F , or P_0 separately. The prover declares a workload of C_{dec} operations. The verifier wants to know if C_{dec} is consistent with P_{obs} . The verifier makes two assumptions:

Assumption 2.1 (Trusted meter). The power trace cannot be forged: P_{obs} differs from the true draw P only by zero-mean noise, which we set to zero throughout.

Assumption 2.2 (Trusted declaration). The declared computation really ran: the smallest total consistent with the window is C_{dec} itself. A power meter cannot by itself confirm this; independently checking that the declared work ran is a separate verification problem—the domain of proof-of-learning and zero-knowledge proof-of-training protocols (Jia et al., 2021; Garg et al., 2023).

Everything else is under the prover’s control and unobserved: the exchange rate $r(t)$, the overhead $P_0(t)$, and how the operation rate $F(t)$ splits between declared and covert work.

2.3. The inverse model

Define G as the forward map that carries a workload to its power trace,

$$G : (F(\cdot), r(\cdot), P_0(\cdot)) \mapsto P(\cdot), \quad (4)$$

defined pointwise by Equation (2). The verifier faces an inverse problem: read the power trace and recover the workload that produced it, i.e. invert G . But G is *non-injective*. Since $r(\cdot)$ is free to vary, a low rate spent on much computation draws exactly the same power as a high rate spent on little. Distinct workloads (F, r, P_0) map to identical traces, and a many-to-one map has no inverse. Inversion therefore returns the identified set $\mathcal{A}(P_{\text{obs}})$ which collects every physically admissible workload that reproduces P_{obs} (Figure 1).

The question “*how many operations ran over the window?*” corresponds to one specific projection of $\mathcal{A}(P_{\text{obs}})$. Define the total compute

$$C = \int_0^T F dt \quad [\text{op}]. \quad (5)$$

Reading C off each admissible workload projects the identified set onto the C -axis,

$$\{C : (F, r, P_0) \in \mathcal{A}(P_{\text{obs}})\} = [C_{\text{lo}}, C_{\text{hi}}]. \quad (6)$$

We define β as the width of that interval, measured against the machine’s peak declared-format capacity $C_{\text{max}} = F_{\text{max}}^{\text{dec}} T$,

$$\beta \equiv \frac{C_{\text{hi}} - C_{\text{lo}}}{C_{\text{max}}}. \quad (7)$$

Under Assumption 2.2 the lower edge is the declaration itself, $C_{\text{lo}} = C_{\text{dec}}$. The interval runs from that honest baseline up to the most compute the meter cannot rule out, so the upper edge is the declared work plus the largest covert load a cheater could slip past the meter,

$$C_{\text{hi}} = C_{\text{dec}} + C_{\text{cov}}^{\text{max}}. \quad (8)$$

Substituting $C_{\text{lo}} = C_{\text{dec}}$ and Equation (8) into Equation (7) leaves a single unknown,

$$\beta = \frac{C_{\text{hi}} - C_{\text{lo}}}{C_{\text{max}}} = \frac{C_{\text{cov}}^{\text{max}}}{C_{\text{max}}}. \quad (9)$$

The width $C_{\text{hi}} - C_{\text{dec}} = C_{\text{cov}}^{\text{max}}$ is exactly the covert operations a cheater could hide. A wider interval means a larger β and more room for undeclared work; $\beta \gg 1$ means the ambiguity exceeds the entire machine.

3. Derivation of β

We now derive a closed form for β . Two physical constraints cap the covert work a window can hold, and each contributes a width. The first is the *energy budget*: the total the meter reports must stay explainable as honest declared work, so covert operations can spend only the joules that ambiguity leaves free (Section 3.1). The second is *capacity*: covert operations must also fit on the machine, in whatever machine-time the declared work leaves unused (Section 3.2). The floor β is the smaller of the two widths, and Section 3.3 evaluates it for the weakest verifier we consider, a bare power meter.

3.1. The energy budget

The derivation of the first width is straightforward: integrate the forward model over the window, split the total energy into declared, covert, and overhead pieces, and ask how large the covert piece can grow before the total stops looking honest. We work with the integrated trace (i.e. the total energy) rather than the resolved time series: treating power as affine in the operation rate (Equation (2)) lets its integral depend on the workload only through time-integrated quantities.

Integrating the forward model Equation (2) over $[0, T]$, the left-hand side is the total energy the meter reports and the right-hand side falls into two pieces,

$$E = \underbrace{\int_0^T r(t) F(t) dt}_{\text{compute energy } E_c} + \underbrace{\int_0^T P_0(t) dt}_{\text{overhead } E_0}. \quad (10)$$

The overhead E_0 is a nuisance the verifier knows only within a band $E_0 \in [E_{0,\text{lo}}, E_{0,\text{hi}}]$. We do not “subtract off” a known baseline: the prover controls fans, clocks, and idle states, so a known-and-removed P_0 would hand it a hiding channel.

The compute energy E_c splits into a declared block and a




Figure 1. Non-injectivity of the forward map. (a) Two workloads (A little compute at a high rate, B much compute at a low rate) whose products rF coincide, so (b) they drive the identical power trace. (c) Inverting recovers only the set of consistent explanations $\mathcal{A}(P_{\text{obs}})$: the rate band $[r_{\text{lo}}, r_{\text{hi}}]$ cuts the constant-energy arc $A-B$, whose projection onto the C -axis spans $[C_{\text{lo}}, C_{\text{hi}}]$; relative to capacity that width is the floor β . The faint grey ribbon is meter noise (ignored in this work). Each rate is drawn constant only for clarity; in reality $r(t)$ varies in time. This does not matter for a question about total operations, which depends on $r(t)$ only through its op-weighted mean $\bar{r} = E_c/C$ (compute energy $E_c = \int_0^T rF dt$ over total operations C). That mean is panel (c)'s vertical axis, so letting $r(t)$ vary leaves the picture unchanged.

covert block,

$$\int_0^T r(t) F(t) dt = \underbrace{\int_0^T r^{\text{dec}}(t) F_{\text{dec}}(t) dt}_{\text{declared}} + \underbrace{\int_0^T r^{\text{cov}}(t) F_{\text{cov}}(t) dt}_{\text{covert}}. \quad (11)$$

The declared block runs $C_{\text{dec}} = \int_0^T F_{\text{dec}} dt$ operations, taken known, and the covert block runs the excess $C_{\text{cov}} = C - C_{\text{dec}}$; this is the compute we are trying to bound. Write each block's energy as its operation count times a single mean cost per operation, i.e. the op-weighted averages

$$\bar{r}^i \equiv \frac{1}{C_i} \int_0^T r^i(t) F_i(t) dt = \int_0^T r^i(t) \frac{F_i(t)}{C_i} dt, \quad (12)$$

for $i \in \{\text{dec}, \text{cov}\}$. The second form makes the averaging explicit: the weight $F_i(t) dt/C_i$ is the fraction of the block's operations that run in $[t, t + dt]$; it is nonnegative and integrates to one. The machine therefore draws

$$E = \underbrace{\bar{r}^{\text{dec}} C_{\text{dec}}}_{\text{declared work}} + \underbrace{\bar{r}^{\text{cov}} C_{\text{cov}}}_{\text{covert work}} + \underbrace{E_0}_{\text{overhead}}, \quad (13)$$

and the verifier sees only E .

The two mean costs are bounded, $\bar{r}^{\text{dec}} \in [r_{\text{lo}}^{\text{dec}}, r_{\text{hi}}^{\text{dec}}]$ and $\bar{r}^{\text{cov}} \geq r_{\text{lo}}^{\text{cov}}$.¹ These band edges,

$$(r_{\text{hi}}^{\text{dec}}, r_{\text{lo}}^{\text{dec}}, r_{\text{lo}}^{\text{cov}}), \quad (14)$$

are the whole parameter vector of this paper: everything below—every rung of the ladder, every measurement, every

¹No upper edge is needed for the covert cost: $\bar{r}^{\text{cov}}$ enters the bound only as a divisor in Equation (16), so the worst case sits at its lower edge and a ceiling would never bind—expensive covert work only hurts the adversary.

headline figure—is the same formula evaluated at a different setting of the triple.

The verifier raises no alarm so long as the measured total is explainable by honest declared work at some admissible rate and overhead. The most an honest run could ever draw is the declared work at the top of its band, on top of the highest overhead,

$$E \leq r_{\text{hi}}^{\text{dec}} C_{\text{dec}} + E_{0,\text{hi}}. \quad (15)$$

The test is one-sided because covert work can only add energy: only a total that runs high is suspicious. Substituting the budget Equation (13) into the no-alarm condition Equation (15) and rearranging,

$$C_{\text{cov}} \leq \frac{(r_{\text{hi}}^{\text{dec}} - \bar{r}^{\text{dec}}) C_{\text{dec}} + (E_{0,\text{hi}} - E_0)}{\bar{r}^{\text{cov}}}. \quad (16)$$

A cheater now pushes every free choice on the right-hand side to its limit:

- *Run the declared work cheaply*, $\bar{r}^{\text{dec}} \rightarrow r_{\text{lo}}^{\text{dec}}$. The verifier cannot tell an expensive declared run from a cheap one, so it must allow the draw all the way up to $r_{\text{hi}}^{\text{dec}} C_{\text{dec}}$; the gap $(r_{\text{hi}}^{\text{dec}} - r_{\text{lo}}^{\text{dec}}) C_{\text{dec}}$ is free energy the prover spends covertly while the total still reads as honest work.
- *Idle the overhead*, $E_0 \rightarrow E_{0,\text{lo}}$, freeing the full band width $\Delta E_0 = E_{0,\text{hi}} - E_{0,\text{lo}}$.
- *Run the covert work cheaply*, $\bar{r}^{\text{cov}} \rightarrow r_{\text{lo}}^{\text{cov}}$, so each hidden operation costs as little as the hardware allows.

At those values Equation (16) attains its maximum, the largest covert load the *energy budget* admits,

$$C_{\text{cov}}^{\text{max, en}} = \frac{(r_{\text{hi}}^{\text{dec}} - r_{\text{lo}}^{\text{dec}}) C_{\text{dec}} + \Delta E_0}{r_{\text{lo}}^{\text{cov}}}. \quad (17)$$




Figure 2. Why the clip carries a γ . The binding resource is the window, not a count of operations: declared work takes hT of the device’s time and the covert block gets the remaining $(1-h)T$ (top). Operations are then whatever each block’s format retires in its own share (bottom), drawn here at $h = 1/2$ and $\gamma = 2$ (int8 covert work against declared fp16), so the covert ticks fall twice as densely as the declared ones. Counting them, the declared block retires 6 operation-units and the covert block 12: the free half window hides a full $C_{\max} = 12$ worth of covert work, not the 6 that segmenting the compute axis would suggest. This is why $C_{\max}$ is a declared-format yardstick rather than a ceiling, and what leaves room for β to exceed one.

Divide by $C_{\max}$ and write the declared *hardware utilisation* $h = C_{\text{dec}}/C_{\max}$ for the fraction of the machine’s declared-format capacity the declared work occupies. What results is the identified-set width of Equation (7) with the admissible set cut by the energy constraint alone — the same normalised width as β , computed as though the no-alarm condition Equation (15) were the only thing standing between the prover and unlimited covert work. We write it β_{en} , the *energy-only width*,

$$\beta_{\text{en}}(h) = h \frac{r_{\text{hi}}^{\text{dec}} - r_{\text{lo}}^{\text{dec}}}{r_{\text{lo}}^{\text{cov}}} + \frac{\Delta E_0}{r_{\text{lo}}^{\text{cov}} C_{\max}}. \quad (18)$$

It is a component of the floor, not yet the floor itself: energy is one constraint on covert work, and Section 3.2 adds the other. Pedagogically, Equation (18) describes (how much honest work there is) $\times$ (how uncertain its per-operation cost is), divided by (how cheaply covert work can run), plus a smaller slack from the overhead band. The first term grows with the declared-band width and with the declared utilisation: more legitimate work offers more cover. The second term is the overhead’s contribution: joules the prover frees by idling overhead that the verifier must still allow for.

Note the asymmetry between the numerator and denominator of Equation (18). Declarations and re-execution can narrow the declared band width $r_{\text{hi}}^{\text{dec}} - r_{\text{lo}}^{\text{dec}}$ in the numerator, but no declaration reaches the covert floor $r_{\text{lo}}^{\text{cov}}$ in the denominator; only hardware-wide observables (clock, voltage) can raise $r_{\text{lo}}^{\text{cov}}$. The two edges therefore also have to be *measured* differently: the declared width tolerates a total-quotient estimate, the covert floor admits only a marginal one. Appendix B states that condition precisely, and it asks less of the hardware than the affine map of Equation (2) does.

3.2. The capacity clip

The second constraint is capacity: covert work must also *fit on the machine*. What binds here is not joules but machine-time. The declared and covert streams multiplex one device across a single window of length T , so the covert block runs only in the time the declared block leaves free. Turning that free time into a count of operations means multiplying by a rate, and the rate on offer is the peak of whichever format the covert stream executes in—which makes this constraint format-dependent from the outset. The setup already bounds the operation rate by the hardware capacity (Section 2.1), but that peak $F_{\max}$ is not one number: it is a property of the executing format. An A100’s dense peaks run from 19.5 Top/s at fp32 on the CUDA cores through 156 at tf32 and 312 at fp16/bf16 to 624 Top/s at int8 on the tensor cores (NVIDIA Corporation, 2020). We write $F_{\max}^{\text{dec}}$ for the peak of the declared format, which fixes the capacity normaliser $C_{\max} = F_{\max}^{\text{dec}} T$ used throughout, and

$$\gamma = \frac{F_{\max}^{\text{cov}}}{F_{\max}^{\text{dec}}} \quad (19)$$

for the capacity ratio of a covert format against the declared one ($\gamma = 2$ for int8 covert work against fp16 declared work on an A100).

However much energy slack the trace leaves, then, the covert block can never retire more than the free window admits. Three steps fix how much that is. Declared work of $C_{\text{dec}} = h C_{\max}$ operations cannot be retired in less than $C_{\text{dec}}/F_{\max}^{\text{dec}} = hT$ of the window, so at most $(1-h)T$ is left free. Run at the covert format’s peak, that free time retires $F_{\max}^{\text{cov}} (1-h)T$ covert operations. Dividing by $C_{\max} = F_{\max}^{\text{dec}} T$ turns the count into a second width, in the same units as β_{en} :

$$C_{\text{cov}}/C_{\max} \leq (1-h)\gamma. \quad (20)$$

The factor γ is the whole content of that clip, and it is why the budget has to be drawn on the time axis rather than the compute axis. It is tempting to segment $[0, C_{\max}]$ directly—declared work takes h of it, covert work the remaining $(1-h)$ —but that treats an operation as a format-independent unit, and it is not: an int8 operation and an fp16 operation consume different shares of the same machine. $C_{\max}$ is a normaliser, the machine’s capacity priced in *declared-format* operations, not a ceiling on how many operations can run. When $\gamma > 1$ the covert stream retires more work per unit of machine-time than the declared format can, so $C_{\text{dec}} + C_{\text{cov}}$ may exceed $C_{\max}$ and the clip is $(1-h)\gamma$ rather than $(1-h)$ (Figure 2). It is also what leaves room for $\beta > 1$: with int8 covert work against declared fp16 ($\gamma = 2$), the free half-window at $h = 1/2$ already retires a full $C_{\max}$ worth of covert operations, and at lower utilisation the clip sits above one, so the width is capped by the energy budget rather than by the silicon.

We evaluate Equation (20) at the *published peak* capacities rather than the rates our kernels achieve: peaks give the covert stream more room than it can really use, so this is the adversary-favorable—and for an upper bound on the ambiguity, conservative—direction (Appendix C quantifies the difference on our measured peaks). Each width is what Equation (7) returns when one constraint is imposed and the other ignored, so the floor—the width when both hold at once—is the smaller of them,

$$\beta(h) = \min[\beta_{\text{en}}(h), (1-h)\gamma]. \quad (21)$$

What kind of bound this is. The maximisation that produced Equation (21) was taken over a *box*: each of the three bands ranges freely, so the prover may sit at $r_{\text{lo}}^{\text{dec}}$, $E_{0,\text{lo}}$ and $r_{\text{lo}}^{\text{cov}}$ all at once. That combination is the *cheating corner*, and nothing in the derivation guarantees a single schedule reaches it.

Assumption 3.1 (Joint feasibility of band corners). There exist admissible workloads attaining $(\bar{r}^{\text{dec}}, E_0, \bar{r}^{\text{cov}}) = (r_{\text{lo}}^{\text{dec}}, E_{0,\text{lo}}, r_{\text{lo}}^{\text{cov}})$ jointly and under shared constraints: the *same* observation window T , the *same* declared work C_{dec} , and one device in a *common thermal and clock state* whose resource-time the two streams multiplex at their *achieved* throughputs.

The assumption is not free, and the overhead band shows why. Our $E_{0,\text{lo}}$ is measured on an idle card (Section 4.2), whereas the window Equation (21) prices is running a declared burst and a covert fill back to back; a window like that does not draw an idle card’s overhead. The bands are also coupled through the clock and the thermal state, which the box treats as independent. So the corner is an outer relaxation of what a device can schedule, and

$$\text{true identified-set width} \leq \beta(h), \quad (22)$$

with equality only under Assumption 3.1. The direction is the safe one for a security claim—we concede the adversary more than the silicon may allow—but β is a worst case in the strong sense, and the gap in Equation (22) is unquantified by the formula alone. Closing it from below takes a constructed schedule rather than an algebraic bound, which is what Section 5.1 does.

One further piece of bookkeeping keeps the box honest. Because both the energy per operation and the capacity are format-dependent, we evaluate the floor inside a single jointly feasible scenario—one declared format and one covert format sharing one window, with the declared band, C_{max} and h all referring to the declared format, and the covert cost and covert peak to the covert format. Combining extrema that no single schedule attains, such as pricing one format’s dearest cells against another’s capacity, is excluded




Figure 3. The shape of the floor. The energy-only width $\beta_{\text{en}}(h) = s h + b$ rises with declared utilisation—more honest work offers more cover (dashed)—while the capacity clip $(1-h)\gamma$ falls, since more honest work leaves less machine free (dotted). The floor Equation (21) is the lower of the two (solid): a tent peaking at the crossing (h^*, β^*) of Equation (23), with energy binding to its left, capacity to its right, and $\beta(1) = 0$ because a fully occupied machine leaves nowhere to hide. No numeric scale is intended—the s , b and γ drawn here are illustrative, and Figure 8 plots the same construction on measured bands.

by construction. The prover may switch formats between the two streams at no cost, which is again the adversary-favorable direction.

The two pull in opposite directions (Figure 3). Energy *rises* with h —more declared work offers more cover—while capacity *falls*—more declared work leaves less machine to hide on. At full declared utilisation the machine is spoken for and $\beta(1) = 0$; at idle there is capacity but almost no cover. The worst case sits at the crossing,

$$h^* = \frac{\gamma - b}{s + \gamma}, \quad \beta^* = \gamma \frac{s + b}{s + \gamma}, \quad (23)$$

writing $\beta_{\text{en}}(h) = s h + b$ for its slope and intercept. The crossing form holds when $b < \gamma$; if the overhead term alone exceeds the clip ($b \geq \gamma$) the clip binds everywhere and the maximum sits at the boundary, $(h^*, \beta^*) = (0, \gamma)$. Each side of the minimum is the binding constraint in its own regime: energy binds for $h < h^*$ (the trace itself caps the covert load), capacity binds for $h > h^*$, and both bind at the crossing. Where the crossing sits is an empirical question about s . On jointly consistent bands—one declared format’s own band over the covert format’s marginal cost—the measured slope is of order one (Section 4.2), so the crossing falls at intermediate utilisation and the worst case is a fraction of γ .

The two are not symmetric, though, despite entering Equation (21) as a bare minimum. Only the energy width *creates* ambiguity: pin the declared band and the overhead ($r_{\text{hi}}^{\text{dec}} = r_{\text{lo}}^{\text{dec}}$, $\Delta E_0 = 0$) and β_{en} vanishes while the clip still stands at $(1-h)\gamma$, yet $\beta = 0$ —free capacity hides nothing once the meter converts joules straight back into operations. The clip never adds width; it only caps the width that blindness has already opened. That asymmetry is why we quote β_{en} on its own as a diagnostic— $\beta_{\text{en}} > 1$ says the

energy budget alone cannot rule out a machine’s worth of covert work—but only β is the width of the identified set.

3.3. The power-only floor

A power meter reports total energy and nothing more: to it an operation is an operation, declared or covert, so there is no way to constrain the declared rate separately from the covert one. The bare meter needs no new formula, then—only the widest setting of the triple Equation (14). The declared edges $(r_{hi}^{dec}, r_{lo}^{dec})$ open to the dearest and cheapest per-operation costs the silicon reaches at all (five tested formats in Section 4.2; the hardware’s true range, holding unmeasured int4 and binary paths, is wider still), and r_{lo}^{cov} stays exactly where it is, because no channel ever reached it. We write $\beta_{en}^\emptyset$ for β_{en} at that setting, the superscript $\emptyset$ denoting the empty set of extra channels: the meter is all the verifier has.

This is the weakest verifier we consider and so the most tempting place to merge the two floors into one r_{lo} , reading the first term of Equation (18) as a band ratio $r_{hi}/r_{lo} - 1$. Beyond the estimand objection of Section 3.1—on the A100 the two numbers differ by half again (Table 1)—that form would hide the asymmetry the rest of the paper turns on: every rung of Section 5 narrows the numerator and none touches the denominator, which a single ratio cannot express.

What does change at this rung is the clip. It is laxer here than the $(1 - h)\gamma$ of Equation (20), because a bare meter pins no format on the *declared* side either: the cheating schedule may run the declared operations themselves in the cheap covert format, where they occupy only h/γ of the machine’s resource-time rather than h . Redoing the two-step count of Section 3.2 with that smaller declared share leaves $(1 - h/\gamma)T$ free, worth $\gamma(1 - h/\gamma) = \gamma - h$ machines of covert work. The power-only floor is therefore

$$\beta(h) = \min[\beta_{en}^\emptyset(h), \gamma - h], \quad (24)$$

and this rung is the deliberate exception to the single-scenario discipline of Section 3.2: both declared edges open to the cross-format envelope while C_{max} and h stay the declared-format yardstick. The combination still passes the single-schedule test, since the honest acceptance ceiling is the fp32 rate and fp32 can sustain the declared rate it implies.² We flag the rung as a cross-format envelope by intent wherever it is quoted; the two clips coincide at $\gamma = 1$.

Two limits bracket the first term of Equation (18). First is *ambiguity*: if the declared rate is known exactly ($r_{hi}^{dec} = r_{lo}^{dec}$) the term vanishes, since a pinned energy budget converts covert joules straight back into operations. That limit

²At the bare rung’s worst case this is $h^*C_{max} = 14.3\text{ Top/s}$ against a measured 17.7 Top/s fp32 peak (Section 4.2).

asks nothing of the covert floor, which is why it is reachable at all: the verifier narrows what it can see and the adversary keeps what it always had. Second is *opportunity*: if the machine is idle ($h = 0$) it vanishes too, with no declared work to hide beneath. The overhead term, by contrast, needs no declared cover at all.

A worked example. To put a number on the power-only floor, take $C_{max} = 100$ operation-units, declared work $C_{dec} = 50$, a rate known only to lie in $[1.5, 2.5]$, and an overhead band $E_0 \in [10, 20]$. In the triple Equation (14) that is $(r_{hi}^{dec}, r_{lo}^{dec}, r_{lo}^{cov}) = (2.5, 1.5, 1.5)$: the example stipulates the two floors equal purely to keep the arithmetic readable, and on real hardware the third entry sits below the second, which raises both terms (Section 4.2). A cheater runs $C = 90$ at the efficient end $\bar{r} = 1.5$ with the overhead idled to $E_0 = 10$, drawing $E = 1.5 \times 90 + 10 = 145$: exactly the energy of honest declared work at the expensive end under high overhead, $2.5 \times 50 + 20 = 145$. No alarm. The verifier cannot separate the honest 50 from the true 90—not for want of statistical power but structurally: the two hypotheses are energy-degenerate, and no test on the total can separate them. Of the 40 covert units, ≈ 33 hid under the rate ambiguity of the declared work (the first term of Equation (18)) and ≈ 7 under the overhead band (the second). In the language of Equation (7): $C_{lo} = 50$, $C_{hi} = 90$, so $\beta = 40/100 = 0.4$. Equation (18) assembles the same number term by term: $\beta_{en}^\emptyset(0.5) = 0.5(2.5 - 1.5)/1.5 + 10/(1.5 \times 100) = 1/3 + 1/15 = 0.4$. The cheat is also capacity-feasible—the 40 covert units fit inside the $(\gamma - h)C_{max} = 50$ units the declared work leaves free (taking $\gamma = 1$)—so here the clip of Equation (24) is slack and $\beta(0.5) = 0.4$. Had the slack exceeded 50 units, the identified set would have ended at the capacity edge instead: at full declared utilisation β vanishes no matter how wide the rate band is.

4. How big is β ?

So far the floor is only a pair of formulas—the energy-only width Equation (18) and its capacity clip Equation (20); this section asks how large they actually are. Because both read the same three numbers, that reduces to one concrete question: how big is each entry of the triple Equation (14)? We first ask what physically drives each entry (Section 4.1), then measure them directly on an NVIDIA A100, where each experiment is built to pin one entry (Section 4.2).

4.1. The two terms

The energy-only width Equation (18) is a sum of two terms: an ambiguity term sh with slope $s = (r_{hi}^{dec} - r_{lo}^{dec})/r_{lo}^{cov}$, set by the spread of the declared exchange rate against the price of a covert operation, and an overhead term

$b = \Delta E_0 / (r_{lo}^{cov} C_{max})$, set by the width of the overhead band against the same price. Both are dimensionless, and both divide by the covert floor: it is the exchange rate at which every slack joule, wherever it comes from, is converted into hidden operations. We take each term in turn.

The exchange-rate term. Its size is governed by the declared band’s width relative to the covert floor, not by any single band’s edge ratio—the two coincide only in the special case $r_{lo}^{dec} = r_{lo}^{cov}$, where s reduces to a ratio minus one. All three entries of the triple are values of the same physical quantity, the energy per operation, so one piece of physics sets the whole vector. A chip’s power divides into a dynamic (switching) part and a static part; the exchange rate comes from the dynamic part, the CMOS switching power $P_{dyn} = \alpha CV^2 f$ (Chandrakasan et al., 1992), whose cost per operation factors into

$$r \approx \kappa(\text{precision}) \cdot \alpha(\text{operands}) \cdot C_{eff}(\text{locality, kernel}) \cdot V^2. \quad (25)$$

Appendix A derives this factorisation; the static part—leakage, cooling, and idle draw—does not enter r but resurfaces as the overhead term below. Each factor is set by a different physical scale of the machine (Figure 4). The numeric-format scale κ falls with datapath width at each precision step (fp32 $\rightarrow$ int8), tracking the pJ-scale energy hierarchy of arithmetic (Horowitz, 2014); the activity factor α is the data-dependent fraction of transistors that toggle for the given operands; the effective capacitance C_{eff} depends on where operands reside and which kernel runs. Less-local operands activate more of the memory hierarchy—storage arrays and peripheral logic, buffers and interfaces, as well as interconnect—so their combined dynamic cost is represented by a larger effective switched capacitance (Horowitz, 2014; Chatterjee et al., 2017); and V^2 is the capacitor-charging energy of the supply voltage, which DVFS slides along the voltage–frequency curve (Pallipadi & Starikovskiy, 2006). With the power meter as the only channel, every one of these factors is unknown to the verifier and conceded at once.

The declared band factorises. That last sentence is the hinge of the paper, so it is worth writing down formally. If each factor of Equation (25) ranges over an interval the verifier cannot narrow, the declared band’s edges are the products of the factor edges,

$$r_{hi}^{dec} = \kappa_{hi} \alpha_{hi} C_{hi} V_{hi}, \quad r_{lo}^{dec} = \kappa_{lo} \alpha_{lo} C_{lo} V_{lo}, \quad (26)$$

writing C for C_{eff} and V for V^2 to keep the subscripts readable. Dividing, the band’s *ratio* is the product of per-factor *widths*

$$\frac{r_{hi}^{dec}}{r_{lo}^{dec}} = w_{\kappa} w_{\alpha} w_C w_V, \quad w_x \equiv \frac{x_{hi}}{x_{lo}}, \quad (27)$$

each $w_x \geq 1$.

Note what Equation (27) returns and what the floor asks for: the factorisation gives the band’s *ratio*, while the slope s of Equation (18) reads its *difference*. Writing $W \equiv w_{\kappa} w_{\alpha} w_C w_V$ for the product, one absolute edge bridges them,

$$s = \frac{r_{hi}^{dec} - r_{lo}^{dec}}{r_{lo}^{cov}} = \frac{r_{lo}^{dec}}{r_{lo}^{cov}} (W - 1). \quad (28)$$

The widths are dimensionless and the slope is not. W fixes how many times over its own floor the band’s ceiling sits; the anchor r_{lo}^{dec} turns that into joules, and only then does the covert floor price it. Setting $r_{lo}^{dec} = r_{lo}^{cov}$ collapses the prefactor to one and recovers the ratio-minus-one reading flagged above—which is why that reading is the special case and not the general one.

With that said, Equation (27) is the map between the two halves of the paper. *This* section measures the four widths, one experiment per factor (Table 1). *Section 5* then observes that a verifier capability which pins a factor sets that factor’s width to 1 and leaves the others untouched: forcing the prover to declare and prove the numeric format sets $w_{\kappa} = 1$ and says nothing whatever about w_{α} , because the operands are not the format. Every rung of the ladder is one such move.

Three cautions on Equation (26). First, it is a worst-case composition—the supremum of a product is the product of the suprema only if the factors can be extreme together, which is the same rectangle assumption the floor already makes of its three bands (Assumption 3.1)—so it is an outer bound. Second, and by Equation (28), a product of widths is not yet a slope: it prices nothing until an absolute edge is supplied alongside it. For both reasons we do not use the factorisation to *evaluate* β . Every number we report comes from measured band edges directly.

The third caution is what that anchor does to the ladder, and it is worth stating before Section 5 leans on it. A ratio is invariant to the *level* a pinned factor is pinned at; a difference is not. Pinning the precision sets $w_{\kappa} \rightarrow 1$ and leaves the residual ratio $w_{\alpha} w_C \approx 1.96$ whichever format is pinned—but composing that residual onto an int8 anchor gives a band [0.786, 1.54] pJ/op and a slope $s \approx 1.5$, and onto an fp32 anchor a band [9.95, 19.5] and $s \approx 19$. Same width removed, same width surviving, slopes differing by the factor just pinned. So “this rung sets $w_x = 1$ ” fixes what a capability removes but not what it is worth: the rung must also say where the surviving band sits. This is why every priced rung of Table 2 reads measured edges—at the precision rung, fp16’s own [1.10, 2.19] and $s = 2.1$ —rather than a running product.

What Equation (27) is for, then, is explaining the ladder

rather than pricing it, and it earns that role by being checkable: Section 4.2 confirms it reproduces both the full measured span and the fp16 sub-band to within a few per cent.

Of the four, only V^2 resists measurement on our platform, which denies user clock control; we therefore carry $w_V \approx 1.5$ from device physics throughout, and flag it as an assumption wherever it appears.³ The other three are measured directly in Section 4.2.

The overhead term. The overhead P_0 —idle draw, cooling, leakage, and I/O—admits no comparably clean factorisation. The exchange rate, though it too varies with the hardware and its operating point, obeys a single switching law that resolves into a few factors with bounded ranges; the overhead obeys no such law. It is a heterogeneous mix of static leakage, cooling control loops, firmware idle states, and platform I/O, so no single relation fixes the band width ΔE_0 from first principles.

4.2. Measurement on an NVIDIA A100

The question is how big β is, so the measurement plan is organised by what β reads—not by the factors of Equation (25), which do not enter the floor separately. Four numbers go in: the three entries of the triple Equation (14) and the overhead width ΔE_0 . Each experiment below exists to pin one of them. Experiment 1 supplies both declared edges ($r_{hi}^{dec}, r_{lo}^{dec}$), Experiment 2 supplies ΔE_0 , and Experiment 3 supplies the covert floor r_{lo}^{cov} .

Each of those four is worth a different number to different verifiers, so we split the bookkeeping by setting rather than by experiment. Table 1(b) collects the four at the one setting that assumes nothing: no format pinned on either declared edge, covert work generic. That is what a bare power meter concedes, and it is the base case the rest of the paper narrows. Values belonging to a capability rather than to the base case are tabulated with their rung instead (Table 2). Mostly that costs no extra measurement—the same Experiment 1 sweep that gives the widest declared pair gives fp16’s own sub-band as one of its bar groups—but one rung needs a workload no experiment here runs, so its measurement travels with it: the useful-work covert floor is Experiment 4, reported in Section 5.

Experiment 1 does a second job on the way, and it is the one the CMOS factorisation asks for. Dividing dynamic switching power by operation rate cancels the clock frequency and leaves κ , α , C_{eff} and V^2 to set the energy per operation, so running the sweep as a factorial decomposes the span between those two declared edges into the separate widths


³Supply voltages on modern accelerators sit near 1 V (Horowitz, 2014) and DVFS moves them by roughly $\pm 10\%$; energy per operation scales as the ratio squared, $(1.1/0.9)^2 \approx 1.5$.

of Equation (27). Those widths price nothing on their own (Equation (28)), but they are the spine Section 5 climbs. The platform did not expose voltage control, so we could not sweep V^2 ; we record the governor-selected clock as operating-point metadata instead. The static component of the power omitted by Equation (25) returns as the intercept P_0 of the affine power law, which is what Experiment 2 reads over the observation window.

Giving the covert floor an experiment of its own needs a word, since Experiments 1 and 2 have already mapped how wide the exchange rate runs. Mapping a range is not the same as pricing its floor. The covert denominator of Equation (18) is a *marginal* quantity that no sweep of whole busy windows estimates (Appendix B), so Experiment 3—and Experiment 4 after it—intervenes on the covert dose directly rather than reading a quotient off a sweep. Experiment 1’s cheapest cell is a value of r_{lo}^{dec} , not a covert floor, and no rung divides by it. One further sweep, Experiment 5, varies the operand axis alone; it corroborates Experiment 1’s α width from independent data and prices nothing, so it is reported in Appendix F rather than here.

Setup and provenance. Experiments 1–3 and 5 ran sequentially on one NVIDIA A100-SXM4-80GB in a single Slurm allocation. Three measurements ran on further cards of the same SKU and are reported where they are used rather than here: the 2–3% device-to-device spread behind the re-execution rung (Section 5), the useful-operand workloads of Experiment 4 (same section), and the matched-energy witnesses of Section 5.1. The platform denied clock locking and power capping, so DVFS is a recorded axis rather than a controlled one: each window logs the governor-selected SM clock and the bands are stratified into clock strata. Energy is read from the on-die NVML total-energy counter; each measured cell fills a ~ 1.5 s window with back-to-back kernels on pre-allocated buffers and keeps the per-cell median over repeated windows. This is a non-adversarial, device-level calibration: the telemetry is trusted as a laboratory instrument, not as the verifier’s sensor in the adversarial setting of Section 2, and one GPU suffices to estimate the bands and the order of magnitude of β ; extending the calibration to off-chip meters and multi-device systems is natural future work. Appendix E gives the full protocol, the meter tolerance, and the provenance manifest, including which card measured which band.

Experiment 1—the declared band, and three of the four factor widths. A compute-bound matmul (operation count $C = 2MNK$ known analytically; the int8 cells retire integer multiply-adds, counted as nominal operations per the convention of Section 2.1), swept factorially over precision (κ , fp32 down to int8), operand locality (C_{eff} , operands resident in vs. streamed through the L2 at a byte-



$$r = \alpha \cdot \kappa \cdot C_{\text{eff}} \cdot V^2$$

f cancels from the marginal rate; leakage and cooling leave as the overhead P_0 , not r

Figure 4. The exchange rate, from transistors to the powered die. Each physical scale contributes one factor to the per-operation energy of Equation (25). (1) A complementary pMOS–nMOS inverter charges or discharges its effective load c_i when its input changes; an operand determines what fraction α of all such cells toggle. (2) Networks of these cells implement the multiply and add of an FMA. The numeric format selects the word width (and supported packing), changing the logic engaged per arithmetic result and hence κ . (3) The FMA reads operands through the register–cache–DRAM hierarchy and writes its result back. Traversing more of that hierarchy switches additional array and peripheral circuits, buffers and interfaces, and interconnect; their combined dynamic cost is represented by C_{eff} . (4) Within a voltage domain, the supply places V across the arithmetic and interconnect loads, so DVFS contributes V^2 per unit capacitance. The clock f cancels from the marginal rate (Equation (34)) and the static draw—leakage, cooling—leaves as the overhead P_0 , not r . Multiplying the factors recovers Equation (25).

identical shape), and operand distribution (α , zeros through adversarial-toggle).

The floor’s two declared entries come off that sweep directly. With no format pinned they are its extremes—the dearest cell 19.5 pJ/op (fp32, dear operands) and the cheapest 0.786 (int8, zeros)—and with the format pinned they are one bar group’s own span, [1.10, 2.19] pJ/op for fp16 (Figure 5). Both pairs are total quotients E/C , the estimand the declared pair admits; the covert floor is not measured here and is never one of these quotients (Table 1).

The sweep is a factorial, though, precisely so that the widths of Equation (27) also come out *separately*: vary one factor with the other three held fixed and the ratio you read off is that factor’s own width, not a pooled spread. Taking each factor’s median conditional ratio over the groups of cells that hold the other two fixed and cover every level of the axis in question,

$$w_\kappa = 13.5, \quad w_\alpha = 1.88, \quad w_C = 1.04, \quad (29)$$

with precision following the textbook ordering $\text{int8} < \text{bf16} \approx \text{fp16} < \text{tf32} < \text{fp32}$. So precision dominates, operand values are worth a factor near two, and locality moves r by four per cent for this compute-bound kernel—a small number, but a measured one, and it is the reason the locality half of the clock/locality rung has so little to buy.⁴

⁴Per access, DRAM costs about a hundred times more than on-chip memory (Horowitz, 2014); the per-operation effect is diluted

Conditioning on the other *designed* axes, rather than on the achieved clock, needs a word of defence, since the clock is exactly the nuisance we cannot lock. Adding it to the conditioning key would fragment a balanced six-level operand axis into groups of one to six levels, and the thin groups contribute ratios near unity that pull the median down—a sparsity artifact, not a cleaner measurement. We therefore condition on the designed axes and *check* the nuisance instead, comparing the median achieved clock across each axis’s levels. The operand and locality levels ran at an identical median clock (balance 1.000 on both), and the precision levels to within 2.7%, so these widths are not clock artefacts. That check is also what separates this experiment from the next one.

The fourth width is missing, and deliberately flagged rather than estimated from this data: the governor held the clock inside 1335–1410 MHz across the whole sweep, a $\pm 3\%$ window, so w_V is *not measured here at all*. The device-physics estimate ≈ 1.5 of Section 4.1 stands in for it, and that gap is why Table 2 declines to price the clock rung.

Two checks that the factorisation is doing real work rather than decorating the argument. Multiplying the three measured widths gives $13.5 \times 1.88 \times 1.04 = 26.5$, against the $24.8\times$ span the pooled cells actually show—so Equa-

by the kernel’s bytes-per-operation, which is why w_C is small here. A memory-bound kernel would not enjoy that dilution, and w_C is the one width we would expect to move substantially with the workload.

quantity	what it is	source	value
<i>(a) the factor widths of Equation (27)</i>			
w_κ	precision	Exp. 1, cond.	$13.5\times$
w_α	operands	Exp. 1, cond.	$1.88\times$
w_C	locality	Exp. 1, cond.	$1.04\times$
w_V	clock (DVFS)	not swept; assumed	$\approx 1.5\times$
$w_\kappa w_\alpha w_C$	product	—	$26.5\times$
<i>(b) the edges the floor reads, nothing pinned</i>			
r_{hi}^{dec}	fp32, dear operands	Exp. 1	19.5
r_{lo}^{dec}	int8, zeros	Exp. 1	0.786
r_{lo}^{cov}	generic covert op	Exp. 3	0.51
ΔP_0	overhead band	Exp. 2	41.2 W

Table 1. What each experiment pins down. Values in pJ/op unless marked otherwise. Block (a) is the physics: the four factors of Equation (25) and how wide each one runs. “Cond.” means the width is measured with the other three factors held fixed, so the widths are separable rather than pooled. Their product reproduces the measured span to 7% (Section 4.2), and $w_\alpha w_C = 1.96$ predicts the fp16 sub-band’s measured 2.00 to 2%—the check that licenses reading Section 5 as one factor pinned per rung. w_V is the exception and is flagged as such: neither Exp. 1 nor Exp. 5 sweeps the clock (the governor held both inside a $\pm 3\%$ window), so its width is a physics estimate, not a measurement. That is why the clock rung of Table 2 is left unpinned. Block (b) is what the floor actually reads: the triple Equation (14) plus the overhead width, at the one setting that assumes nothing—no format pinned on either declared edge, and the covert floor priced for generic covert work. That is the bare meter’s concession and the top rung of Table 2; every capability below it moves one of these four numbers, and Table 2 carries the moved values rather than repeating them here. Nothing in block (a) is an input to block (b): the two meet only through Equation (28), which needs an edge as well as a ratio. Note also that the covert floor is measured differently from the declared edges, and must be—it is the one quantity the floor divides by (Section 3.1 and Appendix B), which is why it sits below the cheapest declared quotient despite describing the same corner of the hardware. C_{max} and γ are read from published peaks rather than measured here.

tion (27) holds to 7%, the residual being factor interaction. And it predicts the effect of a capability before that capability is introduced: pinning the format should set $w_\kappa = 1$ and leave $w_\alpha w_C = 1.96$, while the fp16 band above spans $2.00\times$. Two per cent. Both checks are on ratios, which is the level the factorisation lives at; what pinning the format is *worth* needs that band’s position too, and Equation (28) is where the two meet.

A separate fp16-only sweep corroborates w_α independently, reaching 1.83 where this one reads 1.88; its clock was governor-selected rather than matched, and 1.83 sits below on exactly the side that confound predicts. Appendix F reports it, alongside the construction of the six operand distributions and the resulting ordering of α . Nothing in the floor reads it.

Experiment 2—the overhead band ΔE_0 and the marginal rate. The *rate ladder* is ten fp16 matmul shapes




Figure 5. Experiment 1 supplies both declared entries of the triple Equation (14): the full spread of bars is $(r_{hi}^{dec}, r_{lo}^{dec})$ with no format pinned, and one bar group’s own spread is the pair once that format is declared. Total energy quotient E/C for a fixed matmul kernel on an NVIDIA A100, swept factorially over precision κ (bar groups), operand locality C_{eff} (paired bars), and operand distribution α (whiskers); white points are the individual measured windows ($n = 240$ across 60 cells). Each bar is the median over the operand axis and its whisker that axis’s full range, so the operand factor—a little under $2\times$ inside every format, against locality’s couple of percent—is read off the figure rather than pooled into it. Dashed lines mark the widest setting’s edges, both of them operand extremes: fp32 at its dearest operands, int8 with zeros. Nothing in this figure is the third entry r_{lo}^{cov} , which is a marginal cost and lies below every bar here (Experiment 3).

chosen to span the machine’s sustained throughput: three small cubes (256^3 to 1024^3) that are launch-bound, two skinny- K shapes ($8192 \times 8192 \times 64$ and $\times 256$) that are bandwidth-bound, and five large cubes (2048^3 to 8192^3) that are compute-bound—a $60\text{--}90\times$ spread in sustained operation rate at a single numeric format. Each window runs 100% busy at its shape’s own natural rate. We deliberately do *not* vary F by duty-cycling one kernel against idle: averaging a fixed busy power against a fixed idle power is linear in the duty fraction by construction, so that design would manufacture the affine form rather than test it. Operation counts are analytic ($C = 2MNK$ per matmul, Section 2.1), so the abscissa is an exact op count divided by the measured wall time of the window and the ordinate is metered mean power. Five idle variants—before any CUDA context, context up, context plus operand pool resident, immediately after the hottest ladder cell, and after a 60 s cooldown—supply the $F = 0$ anchors.

The idle range, widened by the adopted $\pm 5\%$ meter tolerance, pins the overhead band to $P_0 \in [59.8, 101.0]$ W (width $\Delta P_0 = 41.2$ W). This band is read directly off those five idle cells, not from a fitted intercept; no input to the floor comes from the fitted slope, which serves only as a diagnostic of the forward model.




Figure 6. The overhead band and the affine forward map (Experiment 2, $n = 65$ windows across a 15-cell rate ladder). *Top*: mean power P vs. nominal-op rate F , plotted as per-cell medians with ± 1 SD error bars over each cell’s repeats. Markers give each cell’s shape class: idle anchors (black, $F = 0$), compute-bound cubes ($M = N = K \geq 2048$, blue), and the launch- and bandwidth-bound mid-ladder shapes (orange). Two affine fits $P = rF + P_0$ are overlaid: the like-for-like fit over the idle anchors and the cubes alone (blue, $r = 1.65$ pJ/op, $R^2 = 0.99$), and the pooled fit over all fifteen cells (black dashed, $r = 1.60$ pJ/op, $R^2 = 0.852$). The slope moves by 4% between them while the scatter grows sharply, so the off-class shapes account for the residual and not for the rate. The shaded band is the observed idle range across five idle variants, $P_0 \in [59.8, 101.0]$ W, read directly off those cells rather than from either fitted intercept. *Bottom*: residuals about the pooled fit, which is what its R^2 measures; the orange shapes carry the scatter. This is a *between-shape* diagnostic of the affine map; the covert denominator the floor actually uses is the *within-cell* marginal cost of Figure 7 (Experiment 3), not either fit here.

As that diagnostic, the affine map holds tightly within a kernel class and loosens across classes. Restricted to the idle anchors and the five compute-bound cubes—like compared with like—the fit gives 1.65 pJ/op at $R^2 = 0.99$ (blue in Figure 6). Opened to all ten shapes it gives 1.60 pJ/op at $R^2 = 0.852$ (black dashed): the slope moves by 4% while the scatter grows sharply, so the launch- and bandwidth-bound shapes account for the residual and not for the rate—the expected behaviour rather than a failure of the model (Appendix B). The derivative the floor actually needs is what Experiment 3 identifies.

Experiment 3—the marginal covert cost. The covert denominator, measured as the theorem defines it: the energy an *added* covert operation costs, with everything else held




Figure 7. The marginal covert cost (Experiment 3): window energy E vs. covert dose C_{cov} for int8 and fp16 covert work, added to a fixed fp16 declared burst inside a fixed 10 s window. *Top*: per-cell medians with ± 1 SD error bars over each cell’s four untrimmed repeats (of five captured; the flagged high-tail repeat is excluded), against the dose fit $E = r_{\text{cov}} C_{\text{cov}} + E_0$; both slopes are linear to $R^2 > 0.999$ (0.53 pJ/op int8, 0.77 pJ/op fp16) and share a zero-dose intercept matching the directly measured zero-dose windows to 0.02%. *Bottom*: residuals. Because the baseline burns identically at every dose, it cancels in the slope: unlike the between-shape fit of Figure 6, this within-cell slope is the incremental estimand $\partial E / \partial C_{\text{cov}}$ the floor’s covert denominator requires (Appendix B).

fixed. Every measured window has the same fixed 10 s wall length and contains the same calibrated fp16 declared burst (utilisation 0.30 of the window, declared-typical operands); to it we add N_{cov} covert matmuls of one format (int8 or fp16, zero operands—the adversary’s cheapest case), then idle-pad to the window. Sweeping the covert dose over six levels from zero to a 0.48 busy fraction (randomised order, drift-reference splices, five repeats per cell) and regressing window energy on covert nominal operations gives the marginal cost as a slope (Figure 7): 0.53 pJ/op for int8 and 0.77 pJ/op for fp16, both fits linear to $R^2 > 0.999$, with the fitted zero-dose intercept agreeing with the directly measured zero-dose windows to 0.02%. Because the fixed window burns the same baseline energy at every dose, baseline power cancels in the slope by construction—the marginal cost sits 0.26 (int8) to 0.39 pJ/op (fp16) below the corresponding total quotients, which amortise it in. The floor uses the one-sided 95% lower confidence limits, $r_{10}^{\text{cov}} = 0.51$ pJ/op (int8) and 0.76 pJ/op (fp16): the covert cost divides the slack, so underpricing covert work is the security-safe direction.⁵ Clock

⁵One-sided in the statistical sense only: the limit controls sampling uncertainty in the fitted dose slope for this workload family,

locking is denied on this platform, so the governor is a recorded, not controlled, axis here too; the window-median clock varies with the dose through the idle-pad share, and the identification rests on the dose-response linearity and the intercept anchor rather than on a common clock stratum—a stated limitation of the platform.

Assembling the floor. The floor is evaluated inside a single jointly feasible fp16-declared scenario (Section 3.2): the declared band is the fp16 cells’ own $[r_{lo}^{dec}, r_{hi}^{dec}] = [1.10, 2.19]$ pJ/op (Experiment 1), the covert cost is Experiment 3’s marginal lower limit, the overhead width is $\Delta P_0 = 41.2$ W (Experiment 2), and capacity and γ (Equation (19)) come from the two formats’ published peaks. The set is format-consistent by construction—our analysis code rejects any cross-format combination of extrema—and joint attainability of its cheating corner within one observation window is Assumption 3.1, which Section 5.1 probes directly. Reading the fp16 cells as the declared band carries one further disclosed assumption: that the declared work is executed in its declared format. No off-site check enforces this; dropping it removes the format pin entirely, so both declared edges open to the full measured envelope and the worst case rises to 1.95 (the bare-meter rung of Section 5). With covert int8 ($\gamma = 2$, $r_{lo}^{cov} = 0.51$ pJ/op) the energy-only width at full declared utilisation is

$$\beta_{en}(1) = 2.39, \quad (30)$$

or 1.62 with covert work held to fp16 ($r_{lo}^{cov} = 0.76$ pJ/op). $\beta_{en}(1)$ is an energy-only diagnostic, not a physical floor: at full declared utilisation the capacity clip $(1 - h)\gamma$ is zero, so no covert work fits and $\beta(1) = 0$. The floor couples the slack to that clip through Equation (21). Figure 8 traces $\beta(h)$ for the two covert scenarios: with covert int8 against declared fp16 the worst case over utilisation is

$$\beta^* = 1.16 \quad \text{at} \quad h^* = 0.42, \quad (31)$$

and with covert work held to fp16 ($\gamma = 1$) it is 0.66 at $h^* = 0.34$. Energy is the binding constraint below the crossing, capacity above it, and both bind at h^* : the worst case sits at intermediate utilisation, where the declared work offers real cover but still leaves much of the machine free.

Both figures are model bounds in the sense of Equation (22): they price the cheating corner, and whether a schedule reaches it—or even reaches the utilisation h^* at which the device, and driver state. It is not a uniform physical lower bound over every execution the threat model admits—unseen kernels, schedules, operand encodings, or thermal states could in principle sit below it. Every rung that divides by r_{lo}^{cov} is therefore conditional on the measured covert families bracketing the adversary’s cheapest edge; the same reading applies to the useful band of Experiment 4 (Section 5), estimated from one trained transformer layer at two shapes.




Figure 8. The capacity-clipped floor $\beta(h)$ (Equation (21)) from the jointly feasible fp16-declared bands, for covert int8 ($\gamma = 2$, blue) and covert fp16 ($\gamma = 1$, orange), both priced at Experiment 3’s marginal lower limits. Faint dashed lines are the unclipped energy-only widths $\beta_{en}(h)$ ($\beta_{en}(1) = 2.39$ and 1.62): energy is the binding constraint left of each crossing, capacity right of it, and both bind at h^* . Dotted lines are the capacity clips $(1 - h)\gamma$; markers are the worst cases of Equation (23).

peak sits—is not something the formula settles. Section 5.1 takes that up.

5. The capability ladder

The bare meter is the worst case: the verifier sees nothing but the power, and the prover has maximum wiggle room. Each capability the verifier gains moves one band edge and shrinks the identified set by a measurable factor.

Section 4 left the ladder with a spine to hang on. The declared band’s ratio is the product $w_\kappa w_\alpha w_C w_V$ of Equation (27), and Section 4.2 measured three of those four widths separately. A verifier capability that pins one factor sets that factor’s width to 1 and—this is the part worth insisting on—leaves the rest of the product exactly where it was. Forcing the prover to declare and prove the numeric format sets $w_\kappa = 1$; it says nothing about which operands ran, so w_α survives untouched. Reading the rungs below as a running product is the quickest way to see both what each buys and what it conspicuously does not.

The product is the spine and not the price, though, and Equation (28) is why: a capability that pins a factor also fixes where the surviving band *sits*, and the same width removed at a different level is worth a different number of joules. Each rung below therefore quotes the measured edges its capability leaves standing, and β^* is recomputed from those edges rather than accumulated down the product. What the running product tells us is which lever is large, and which nuisance the next rung has failed to touch.




Figure 9. The capability ladder of Table 2 on a log scale; every plotted value is the clipped worst case $\beta^* = \max_h \beta(h)$. Filled circles are verifier capabilities (types D, C); the square rung is forced by the threat model (type M: hidden zero-ops are no threat); the hollow rung is an assumption about the adversary, not a capability (type A). The shaded slot is the unpriced prospective clock/locality rung; daggered labels carry the disclosed conditions of Table 2. The $\beta^* = 1$ line marks one whole machine of conceded ambiguity; the top two rungs are model upper bounds whose maximisers lie above the feasible window, so where they fall relative to that line is not something we claim to settle.

Three of the rungs are not of that kind, and the taxonomy is clearer for saying so up front. Re-executing the declared work does not pin a factor: it measures r_{hi}^{dec} directly and so *bypasses* the factorisation altogether, collapsing the numerator to a device-transfer gap no matter which widths were still open—which is why it is the largest single step in the table and why it is the rung that finally prices w_α . And the bottom two rungs act on the denominator, where no capability reaches: only a restriction on what the adversary wants moves r_{lo}^{cov} .

Table 2 and Figure 9 walk down that ladder *cumulatively* — each rung adds one capability to all the rungs above it — and their captions fix the accounting every priced rung shares: the quantity re-computed at each step, the joint model it is computed in, and the two disclosed conditions that thread the chain. Because a verifier that gains a capability can only shrink the set, the priced chain is monotone by construction. The rest of this section takes the rungs in turn.

The bare meter ($\beta^* = 1.95$). With power alone the verifier cannot tell declared work from covert, nor one numeric format from another, so the declared band opens to the full measured envelope of Section 3.3:⁶ the triple takes its widest

⁶“Measured” scopes the envelope to the five formats of Section 4.2 (fp32, tf32, fp16, bf16, int8). The A100 also exposes

setting, $(r_{hi}^{dec}, r_{lo}^{dec}, r_{lo}^{cov}) = (19.5, 0.786, 0.51)$ pJ/op: the declared edges are fp32 at its dearest operands and int8 zeros, and the covert entry is where it always is, at the cheapest incremental operation the silicon admits. No precision is assumed on either declared edge — reading the ceiling as the fp16 cell while opening the floor to int8 would assume the declared format on one edge and deny it on the other, which a bare meter cannot do. The slope is then $s = (19.5 - 0.786)/0.51 = 36$: larger than the $24.8\times$ declared span, not because the band has been stretched but because its width is priced against a covert operation cheaper than either edge of it — large enough that the capacity clip binds even at low utilisation. The clip is the cross-format $\gamma - h$ of Equation (24): with no format pinned, the declared operations themselves execute in the cheap covert format and occupy h/γ of resource-time, not h . The clipped worst case peaks at $\beta^* = 1.95$ near $h^* = 0.046$, essentially the int8 density advantage $\gamma = 2$ that a near-idle machine leaves free. This maximiser sits *inside* the feasible window ($h^* < h_{max} \approx 0.24$), so — unlike the precision rung, whose peak is pushed above the busy-time ceiling — confining the worst case to feasible utilisation does not lower it; the bare-meter figure rests on the corner assumption Assumption 3.1, alongside the measured-scope and covert-edge conditions already flagged. Finally, note that each entry here carries the estimand its role requires (Appendix B): the declared pair is a difference of total quotients, the covert entry a marginal lower bound. One conservatism does survive that bookkeeping—a total-quotient ceiling amortises baseline power into r_{hi}^{dec} while the overhead band is carried again as its own term, so leakage and cooling are charged twice—which widens the rung rather than narrowing it, the safe direction for a worst case (Appendix C quantifies it).

Declare and check the precision ($1.95 \rightarrow 1.16$, conditional). The prover can be made to declare the declared work’s numeric format, and the verifier can re-run that work on its trusted cluster and confirm the outputs are what fp16 execution produces. In the taxonomy of Table 2 this rung is a declaration plus an offline check: nothing new is sensed at the facility — and that is exactly its limit. The check makes the declaration *precise*; it cannot make it *enforced*. In the language of Equation (27) this rung sets $w_\kappa \rightarrow 1$ and nothing else: the band ratio falls from $w_\kappa w_\alpha w_C w_V$ to $w_\alpha w_C w_V$, which is why the fp16 cells measure a residual $2.00\times$ against the 1.96 that $w_\alpha w_C$ predicts (Section 4.2). Knowing the format tells the verifier nothing about which

dense int4 (1,248 Top/s) and binary (4,992 Top/s) tensor-core paths (NVIDIA Corporation, 2020) that no committed measurement prices: an int4 covert edge could sit below the int8 marginal edge and carries $\gamma = 4$ against fp16. Unmeasured formats can only widen this rung, so 1.95 is the model worst case within the measured scope and a lower bound on the true full-hardware ambiguity, not a full-hardware worst case.

verifier gains	type	factor acted on, and the edges it leaves	β^*	$\beta_{\text{en}}(1)$
nothing (bare meter)	—	<i>all open</i> : (19.5, 0.786, 0.51)	1.95	37
precision declared & checked [†]	D, C	$w_\kappa \rightarrow 1$; band $\rightarrow$ [1.10, 2.19]	1.16	2.4
clock & locality observed	T	$w_V, w_C \rightarrow 1$ (<i>unpriced</i>)	—	—
declared work re-executed [‡]	C	<i>bypasses</i> w_α : band $\rightarrow$ gap	0.34	0.35
covert work on useful data [‡]	M	denominator: 0.51 $\rightarrow$ 2.53	0.071	0.072
covert at top of useful band [‡]	A	denominator: $\rightarrow$ 3.16	0.057	0.058

Table 2. The capability ladder. Each rung grants the verifier one more capability (cumulative, top to bottom) and is re-priced as the clipped physical worst case $\beta^* = \max_h \beta(h)$ (Equation (23)) under the single jointly feasible scenario of Section 3.2 (fp16-declared case study, int8 covert edge at the Exp. 3 marginal LCB, $\gamma = 2$; the format-unpinned bare rung carries the cross-format clip $\gamma - h$ of Equation (24)); $\beta_{\text{en}}(1)$ is the energy-only diagnostic, not a floor. Numeric band edges in pJ/op. The third column carries the edges as well as the factor, because a rung that pins a factor fixes where the surviving band sits and only the two together price it (Equation (28)); each β^* is recomputed from the edges on its own row, not accumulated down the ladder. The top rung is Table 1(b) verbatim, the setting that assumes nothing. Its successors move one entry at a time and are measured no differently: the fp16 band [1.10, 2.19] is Exp. 1’s fp16 cells, the useful covert edges 2.53 and 3.16 are Exp. 4’s marginal lower limits, and $\Delta P_0 = 41.2$ W (Exp. 2) is common to every row. Disclosed conditions: [†] — assumes the declared work is executed in its declared format (an off-site check cannot enforce this; Section 5); [‡] — additionally assumes re-execution transfers at an observed operating point. Both are inherited by every rung below their first appearance. The type column says what kind of thing each rung is: D — a declaration by the prover, checked off-site; C — offline characterisation on the verifier’s own hardware, transferring to the facility only up to a stated gap; T — a trusted on-site observable, here left *unpriced* on two independent grounds: attested telemetry does not yet exist, and w_V —the width such a channel would remove—is not measured by any sweep of ours; M — a threat-model restriction (hidden zero-ops are no threat); A — an assumption about the adversary’s objective, not a capability.

operands ran, so the whole activity factor survives—a point the next two rungs exist to address. Concretely the declared pair narrows to the fp16 band’s own [1.10, 2.19] pJ/op ($r_{\text{hi}}^{\text{dec}}$ down from 19.5, $r_{\text{lo}}^{\text{dec}}$ up from 0.786, and $r_{\text{lo}}^{\text{cov}}$ untouched at 0.51), and it buys that only under the disclosed condition that the declared work is actually executed in its declared format at the monitored facility ([†] in Table 2); facility-local format enforcement is an open capability, not one this rung grants. Under that condition the rung is precisely the headline case: $\beta^* = 1.16$ at $h^* = 0.42$ (Equation (31)). Here and below, “the verifier knows the precision” always means the *declared* precision. The covert work’s precision is unobservable—declaring it would make it not covert—so it never enters the numerator; it resurfaces at the bottom rungs.

Observe the clock and locality (unpriced). This rung would set $w_V \rightarrow 1$ and $w_C \rightarrow 1$, leaving the declared band at its activity core w_α alone. Both are hardware-wide observables the power meter does not measure. We deliberately quote no number for it, and the factorisation makes the two reasons crisper than a single “prospective” would.

First, the capability does not exist: every route to it—attested clock telemetry, locked clocks by policy, or the EM side channel (Section 6)—is prospective, so there is no channel to price. Second, and independently, *we have not measured* w_V . Neither Experiment 1 nor Experiment 5 sweeps the clock; the governor held both inside a $\pm 3\%$ window, so the ≈ 1.5 we carry for w_V is Section 4.1’s device-physics estimate and not a measurement (Table 1). Pricing a security guarantee on the width a capability removes re-

quires knowing that width, and for this rung we do not. The sub-factor we *can* measure, $w_C = 1.04$, would on its own buy almost nothing—four per cent of a band whose other factors run to 13.5 and 1.88. So the rung stays in the ladder as the slot a trusted on-site observable would fill, and the honest reading is that its value is unknown rather than small — and the next rung depends on it.

Re-execute the declared work (1.16 $\rightarrow$ 0.34, conditional).

This is the rung that finally prices the activity factor, and it does so by leaving the factorisation behind. Re-executing to *measure* the declared work’s actual energy per operation, not merely confirm its format, reads the operands that actually ran; the prover’s freedom to claim cheap data cancels, so w_α is disposed of without ever being pinned. The same stroke disposes of w_C and w_V : a direct measurement of $\bar{r}^{\text{dec}}$ does not care which factors were still open. So the step is not $w_\alpha \rightarrow 1$ but the collapse of the whole product, $r_{\text{hi}}^{\text{dec}} \rightarrow r_{\text{lo}}^{\text{dec}}$, and the numerator falls to a residual set instead by the *hardware gap*: re-execution runs on the verifier’s trusted cluster, not the monitored facility, and energy per operation is hardware-specific—different accelerators, firmware, cooling, operating point—so the measured r^{dec} transfers only up to a gap. That gap, not zero, is the floor on how far re-execution can pin the declared band. We measured the easiest flavour of it directly: re-measuring every band edge on three further A100s of the same SKU, each agrees to 2–3% across devices (Appendix E), so same-SKU re-execution transfers bands at the few-per-cent level. That is a lower bound on the difficulty; across different accelerators, firmware, and cooling the gap is larger and unmeasured. We take the residual gap at a few per cent of the declared

rate, ≈ 0.05 pJ/op (an assumption informed by the same-SKU spread, not a cross-hardware measurement), giving $\beta^* = 0.34$ at $h^* = 0.83$ — at this depth the overhead band, not the declared residual, carries most of the width. Read the rung as a sensitivity analysis in the gap, not a measured bound: at the same-SKU spread itself (gap 0.02–0.04 pJ/op) the rung reads 0.29–0.32; at our stated 0.05 pJ/op, 0.34; a cross-vendor gap of 0.2 pJ/op would put it at ≈ 0.54 . Measuring that gap — the same declared workload characterised across accelerator generations and vendors — is the experiment this rung waits on.

The rung also inherits a second condition (§ in Table 2): *re-execution only pays at an observed operating point*. Re-execution measures the declared work’s energy per operation at the *verifier’s* operating point; if the prover’s clock state is unobserved — and the previous rung is exactly the one still unpriced — the prover may run the declared work anywhere on the DVFS curve, and the transfer error is not the hardware gap but the full DVFS factor of ~ 1.5 . A verifier with only “power and paperwork” (the power meter plus declarations and re-execution, but no on-site sensing) therefore stalls at $\beta^* \approx 1$: re-executing the declared work buys almost nothing that checking its precision had not already bought.

Restrict covert work to useful data (0.34 $\rightarrow$ 0.071; **Experiment 4**). This rung is forced by the threat model rather than granted by a channel, and the logic runs in steps. Covert work is unseeable *by definition*: no declaration reaches it, so its floor r_{lo}^{cov} is bounded only by the silicon — the int8 marginal edge, 0.51 pJ/op. But the cheapest operations on that edge are operand zeros, and zeros compute nothing useful: hidden zero-ops are no threat.

This rung therefore comes with a measurement of its own, and it is the one experiment Section 4.2 does not report, because no number above this rung reads it. We measured the *incremental* cost of useful covert work directly: an int8-quantised transformer block built from one trained Pythia-6.9B layer, run on real text activations (not synthetic patterns), with the quantisation scales applied and every metered operation — the int8 matmuls and the fp16 attention alike — counted. Sweeping that block as the covert dose inside Experiment 3’s construction — a fixed declared burst inside a fixed window, so baseline power cancels in the dose slope (Appendix B) — at two shapes that bracket the attention share of the op count, and regressing window energy on covert nominal operations, prices genuinely useful covert work at a *marginal* 2.54–3.21 pJ/op, both fits linear to $R^2 > 0.999$. We carry the one-sided 95% lower limits, $r_{useful} = [2.53, 3.16]$ pJ/op, for the reason Experiment 3 gives: under-pricing covert work is the security-safe direction. The standalone total quotient E/C is dearer, [3.18, 4.22] (one-sided lower limits; point means

[3.26, 4.24]), only because it amortises idle power the covert work does not draw. Restricting the adversary to work worth doing raises the covert edge to the bottom of that band, 0.51 $\rightarrow$ 2.53 pJ/op, while the declared-side residual stays at the hardware gap — the covert band’s own width prices nothing on the declared side — and the rung drops to $\beta^* = 0.071$: every slack joule now buys the adversary $\approx 4.9\times$ fewer useful operations than silicon-floor accounting would concede. Because this edge is an incremental dose slope, not a standalone total quotient, it is a genuine marginal lower bound on the covert cost — the estimand the denominator requires (Appendix B), measured the same way as Experiment 3’s int8 covert edge.

Covert work at the top of the useful band (0.071 $\rightarrow$ 0.057). The final rung is an assumption, not a capability. Only by assuming the covert work forgoes all of its remaining freedom inside the useful band and pays that band’s ceiling, $r_{lo}^{cov} \rightarrow 3.16$ pJ/op, does the residual fall to ≈ 0.057 . A bound of that order — covert compute at several per cent of capacity — exists only on this rung, conditional on an assumption about the adversary’s objective rather than on any measurement; the distinction matters for the power-monitoring layers proposed in the verification literature (Baker et al., 2025).

5.1. How much of the bound is real?

Every number so far, at every rung, is the same kind of object: a maximum over the band box of Assumption 3.1, and therefore an upper bound on the identified-set width rather than the width itself (Equation (22)). This section asks how much of that headroom a real schedule can occupy. The question is not a rung — no verifier gains anything here — but it applies to every rung, and it separates what the model concedes from what the silicon delivers.

What β is, and is not. Two readings of the width must be kept apart. Given an *observed* total E , the identified set has a width that depends on how much energy was actually drawn. β is instead the largest such width over every total the no-alarm rule accepts: a design-time statement about what *any* accepted window could hide, not the ambiguity of the window in front of you. Quoting the second as though it were the first overstates a typical audit, and quoting the energy-only width β_{en} without its capacity clip overstates what the silicon admits. Read with those caveats, β is a design-time budget: whatever the audited window turns out to be, no accepted window can hide more than this.

The maximiser need not be occupiable. The first gap is on the utilisation axis. We evaluate the capacity clip at *published peak* rates (Section 3.2), which deliberately gives the covert stream more room than a real kernel can use,

so the clip does not know what actually fits in a window. Our witness construction below does: above a busy-time ceiling $h_{\max} \approx 0.24$ its cheating arm’s work no longer fits T . The headline $\beta^* = 1.16$ sits at $h^* = 0.42$, nearly twice that — so at the headline utilisation the wall-clock budget binds before any schedule reaches the corner. Restricting the maximisation to $h \leq h_{\max}$ puts the model worst case at $\beta \approx 0.77$. The ceiling is a property of that construction rather than a general feasibility bound, and a different covert workload would move it (Appendix D); the general point is that a verifier reading *time* as well as energy constrains more than the energy channel alone prices, and that joint energy–time inversion is future work. The bare-meter rung is the instructive contrast: its maximiser $h^* = 0.046$ sits well inside the feasible window, so confining that rung to feasible utilisation does not lower it at all.

Matched-energy witnesses: the bound from below. The second gap is the corner itself, and it admits a direct empirical probe. Matched-pair witnesses on one A100 hold the trusted declaration and the observation window fixed by construction: run A executes N declared fp16 iterations on dear operands (the $r_{\text{hi}}^{\text{dec}}$ edge) and idles to a fixed $T = 10$ s window; run B executes the *same* N iterations on cheap operands (the $r_{\text{lo}}^{\text{dec}}$ edge) plus a covert int8 fill tuned to match A’s energy, padded to the same window. At $h = 0.10$ the two arms match to $|\Delta E|/E = 0.9\%$ while B hides 3.5×10^4 covert int8 iterations, and the one-sided no-alarm rule, set three standard deviations above the honest mean, accepts *every* cheating repeat. The witnesses reach $h = 0.10, 0.15$ and 0.216 , where the accepted covert load reaches $0.19, 0.29$ and 0.41 machines of declared-format capacity respectively — measured, capacity- and time-feasible lower bounds on the ambiguity (full protocol in Appendix D).

They are lower witnesses, not corner attainment. At $h = 0.216$ the accepted load 0.41 sits well below the model’s $\beta(0.216) = 0.72$, and the achieved rates say why: 1.84 pJ/op on the dear declared arm against the band ceiling 2.19 , and 0.54 on the covert fill against the marginal edge 0.51 . The construction approaches every edge of the box and reaches none of them, which is the relaxation gap of Equation (22) measured rather than assumed.

The resulting bracket. Putting the two together, the identified-set width a bare meter admits against a format-checked declaration is *at least* 0.41 machines — measured, attained, and accepted by the sensor — and *at most* about 1.16 as a model bound whose maximiser is infeasible, or about 0.77 if the bound is confined to a feasible window. Four-tenths of a machine of covert compute passing an honest declaration is a wide identified set; whether it crosses the whole-machine “vacuity” line is not something the attained witness settles, and we do not claim it does. The width is set

jointly by the energy budget and the silicon, not by either alone.

6. Discussion

The meter is an anchor, not an instrument. The result reads pessimistic — against a bare power meter a matched-energy witness demonstrably hides four-tenths of a machine of covert compute, and the model worst case runs higher still, to $\beta^* = 1.95$ against a bare meter — a figure that rests on the joint feasibility of the band corners (Assumption 3.1) — but the ladder inverts the reading. The meter’s role in a verification regime is not to count operations by itself; it is to be the one channel whose integrity needs no trust in the prover, the anchor against which declarations, attested clocks, and re-execution are checked. Power plus checked declarations reaches $\beta^* \approx 1.2$; adding an observed operating point and re-execution of the declared work reaches ≈ 0.34 , and restricting covert work to useful data ≈ 0.071 of capacity. Every rung is a protocol requirement a treaty could stipulate, together with a disclosed condition it rests on, and Equation (23) states exactly which band each requirement moves and by how much. The negative result is equally operational: because re-execution transfers only at the verifier’s operating point, “power plus paperwork” — declarations and re-execution with no on-site sensing — stalls at $\beta^* \approx 1$, a regime design worth ruling out before it is proposed.

What no verifier capability reaches. The denominator of Equation (18) is a statement about the adversary, not the instrumentation. Covert work is by definition undeclared, so its exchange rate is bounded only by the silicon and by what the threat model says the adversary wants: excluding zero operands (which compute nothing) raises the covert floor from the int8 marginal edge to the cheapest *useful* rate — a measured $4.9\times$ step, 0.51 to 2.53 pJ/op — and the remaining $1.25\times$, the useful band’s own width, closes only by assuming the adversary declines cheap-but-useful work — sparse activations, quantised weights — that our own measurements show exists. Any bound quoted for power-based verification should state which side of that assumption it stands on.

From one device to a facility. Our bands are device physics, measured on one A100; a facility is a sum of devices behind shared power delivery and cooling, and we deliberately do not offer a fleet-average version of Equation (18). Scaling the floor up is a *multi-resource allocation problem*, not a change of units: each device carries its own capacity clip and utilisation, the prover chooses how declared work is spread across devices (and formats) to maximise cover, and the facility meter adds cooling and conversion overheads that the prover partially controls. The

identified set is then the projection of a per-device product of constraints, and its width depends on the allocation the prover is free to pick; treating h as a single fleet average can under- or overstate it. What does carry over qualitatively: overhead bands widen at facility scale, and cheaper datapaths widen the precision envelope beyond our $24.8\times$ and raise the capacity ratio γ of the cheapest format—int4 and binary tensor-core paths on the A100 itself (unmeasured here; Section 5), FP8 and FP4 on newer parts. Making the multi-device floor precise—including which per-device telemetry restores the single-device analysis—is open.

Limitations. Seven caveats bound the claims. (i) The bands come from one vendor, one GPU model, one driver, and one cluster; the ladder’s numbers are A100-specific, though the closed form is not. (ii) The band edges are read from a factorial sweep whose corner cell (int8, zero operands) is directly measured, but on synthetic operands; genuinely useful covert work prices at the measured marginal useful band $[2.53, 3.16]$ pJ/op instead (Experiment 4’s dose-response; the standalone total-quotient lower limits are $[3.18, 4.22]$), and the bottom rungs depend on that band’s width. (iii) The hardware gap behind the re-execution rung (≈ 0.05 pJ/op) is an assumption. The same-SKU piece is grounded—band edges agree to 2–3% across three A100s (Section 4.2)—but the cross-hardware gap (different accelerators, firmware, cooling) remains unmeasured. (iv) DVFS was never swept — the platform denied clock-locking — so the voltage factor rides on literature numbers rather than our own; persistence mode was likewise fixed on, so its contribution to the controllable idle range went unmeasured. (v) Energy was read from on-die telemetry in a non-adversarial lab characterisation (Section 4.2); the counter carries no published GA100 accuracy specification, so the $\pm 5\%$ tolerance we apply is adopted from NVML’s documented Fermi/Kepler figure, and the ~ 100 ms telemetry cadence is an empirical characterisation (Yang et al., 2023), not a vendor specification; an off-chip meter would add its own calibration band, which only widens β . (vi) On shared nodes, windows stalled by co-tenants were dropped by a pre-specified power guard, and high-tail windows are flagged rather than deleted; all raw rows remain on disk. The flagged rows are excluded from the published bands and regressions—including the fp16 declared band and the marginal-dose fits, not only the descriptive medians—and with five repeats per cell the 20% trim always removes one repeat, excursion or not. The sensitivity is small: restoring every positive flagged record moves the fp16 band to $[1.10, 2.20]$ pJ/op, the int8 marginal lower limit to 0.510, and the headline to $\beta^* = 1.162$. We claim no single statistical coverage guarantee over all the quantities that enter the bound: they are of three different kinds — a coverage-controlled confidence limit, a band width, and conservative design parameters — and Appendix B states the status of

each. (vii) The locality contrast is attenuated in the fp32/tp32 cache-resident cells, whose operands marginally spill L2.

What richer telemetry adds. Everything here is time-integrated, and the sufficiency of total energy for a question about totals is a consequence of our *free-variation* model: because $r(t)$ and $P_0(t)$ may vary arbitrarily within their bands, the trace’s shape carries no information about totals beyond its integral. Real hardware is not perfectly free—power-ramp limits, DVFS governor dynamics, the instantaneous clip $F(t) \leq F_{\max}$, and any *declared schedule* the prover commits to all constrain how the trace can move, and a verifier who models them can extract more from the resolved trace than from its integral. Conditional on the free-variation model, then, the time *profile* adds nothing about how much compute ran; unconditionally, that is an open modelling question, not a theorem. The profile also carries different information — workload cadence, training-step periodicity — that bears on what *kind* of compute ran rather than how much; that question needs different machinery and we do not treat it here.

The same caveat applies beyond the time axis. Our result prices the integrated-energy channel under free variation; a verifier with richer instrumentation is outside the model. Multiple simultaneous channels (per-rail power, electromagnetic, thermal), cryptographically authenticated telemetry (the proof-of-learning and zero-knowledge proof-of-training protocols of Assumption 2.2), or independently controlled measurements the prover cannot shape (locked clocks, tamper-resistant attested counters) each carry constraints the integral does not, and any of them may narrow the identified set below the widths reported here. The impossibility we state is for total-energy sensing under the free-variation threat model, not for analogue sensing in general.

That none of these channels is available yet is what leaves the clock-and-locality rung of Section 5 unpriced. Every route to it is prospective: firmware-reported clocks count only over a tamper-resistant read path that does not yet exist, locked clocks are a policy the prover must accept, and the EM side channel is unproven for this purpose—switching at f radiates at f and its harmonics, and magnetic emanations from a GPU’s power delivery are measurable in practice (Maia et al., 2022), but clock-line radiation itself has not been demonstrated, and the voltage would then have to follow from the hardware’s V – f table. The locality half is more speculative still—per-rail metering (core vs. HBM regulators) or interconnect telemetry might resolve it—though Experiment 1 found locality moves r little for compute-bound work. Pricing that rung is the instrumentation question the ladder leaves open.

7. Conclusion

A bare power meter cannot count operations. The energy cost of an operation is unknown and data-dependent—across the five formats we measured the envelope spans $24.8\times$ on an A100 (0.79–19.5 pJ/op), with the hardware’s unmeasured int4 and binary paths wider still, and even inside one jointly feasible declared format the band and the marginal covert cost leave an energy slack ($\beta_{\text{en}}(1) = 2.39$, Equation (30), an energy-only diagnostic since no covert work fits at full utilisation). Coupling that slack to silicon capacity gives the clipped floor of Equation (21). What we attain and the sensor accepts is a matched-energy witness hiding 0.41 machines of covert int8 work at $h \approx 0.22$; the model worst case peaks higher, $\beta^* = 1.16$ at $h^* = 0.42$ (Equation (31), 1.95 against the bare meter with the declared format unenforced, within the measured five-format scope), but that peak is an upper bound whose maximiser sits above the busy-time ceiling—confined to a feasible window the model worst case for the precision-checked sensor is about 0.77. The identified-set width the meter admits is thus at least 0.41 machines and at most about one, set jointly by the energy budget and the silicon. This is not a statistical shortfall a better meter or a longer window repairs; it is a structural degeneracy of the integrated-energy forward map under free variation — a property of the model, not of the A100. The specific widths above (the $24.8\times$ span, $\beta_{\text{en}}(1) = 2.39$, the measured 0.41, and the model $\beta^* = 1.16/1.95$) are contingent A100 estimates, scoped to the formats we measured; the closed form they instantiate is device-agnostic. An identified set this wide—four-tenths of a machine of covert compute demonstrably passing an honest declaration—substantially weakens the raw joules-to-FLOPs conversion in the governance literature against an adversary free to vary energy intensity.

But the floor is a closed form built from named physical bands, and that is what makes it useful: each verifier capability moves one band edge by a measured amount, and every rung is re-priced as the same clipped worst case $\max_h \beta(h)$. Precision declared and checked brings the measured-scope worst case from 1.95 to 1.16, conditional on the declared work executing in its declared format; re-execution of the declared work at an observed operating point to ≈ 0.34 — a residual that rests on a stated hardware-gap assumption, not a measurement (Table 2). The rungs interlock — re-execution pays only at an observed operating point, so power plus paperwork stalls at ≈ 1 — and no capability reaches the covert side: restricting covert work to the measured useful-operand band brings ≈ 0.071 , and the final collapse to ≈ 0.057 closes only by assumption on the adversary. This paper prices each rung, with its conditions stated, so that a verification regime can be designed — or rejected — on numbers rather than on the intuition that joules imply FLOPs.

References

- Ansari, S. Hardware-level governance of AI compute: A feasibility taxonomy for regulatory compliance and treaty verification, April 2026. URL <https://arxiv.org/abs/2604.04712>.
- Arefin, S. E. and Serwadda, A. Exploiting HDMI and USB ports for GPU side-channel insights, October 2024. URL <https://arxiv.org/abs/2410.02539>. Published at ACM CODASPY 2025.
- Baker, M., Kulp, G., Marks, O., Brundage, M., and Heim, L. Verifying international agreements on AI: Six layers of verification for rules on large-scale AI development and deployment, July 2025. URL <https://arxiv.org/abs/2507.15916>.
- Bengio, Y., Clare, S., Prunkl, C., Andriushchenko, M., Bucknall, B., et al. International AI safety report 2026, February 2026. URL <https://arxiv.org/abs/2602.21012>.
- Chandrakasan, A. P., Sheng, S., and Brodersen, R. W. Low-power CMOS digital design. *IEEE Journal of Solid-State Circuits*, 27(4):473–484, 1992. doi: 10.1109/4.126534.
- Chatterjee, N., O’Connor, M., Lee, D., Johnson, D. R., Keckler, S. W., Rhu, M., and Dally, W. J. Architecting an energy-efficient DRAM system for GPUs. In *2017 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, pp. 73–84, 2017. doi: 10.1109/HPCA.2017.58.
- Chaudhuri, A., Shukla, S., Bhattacharya, S., and Mukhopadhyay, D. “Energon”: Unveiling transformers from GPU power and thermal side-channels, August 2025. URL <https://arxiv.org/abs/2508.01768>.
- Choukse, E., Warriar, B., Heath, S., Belmont, L., Zhao, A., Khan, H. A., Harry, B., Kappel, M., Hewett, R. J., Datta, K., et al. Power stabilization for AI training datacenters, August 2025. URL <https://arxiv.org/abs/2508.14318>. Microsoft.
- Clark, S. S., Ransford, B., Rahmati, A., Guineau, S., Sorber, J., Xu, W., and Fu, K. WattsUpDoc: Power side channels to nonintrusively discover untargeted malware on embedded medical devices. In *2013 USENIX Workshop on Health Information Technologies (HealthTech 13)*. USENIX Association, 2013.
- Gangwal, A., Piazzetta, S. G., Lain, G., and Conti, M. Detecting covert cryptomining using HPC, 2019. URL <https://arxiv.org/abs/1909.00268>. Published at CANS 2020.

- Gao, Y., Qiu, H., Zhang, Z., Wang, B., Ma, H., Abuadbbba, A., Xue, M., Fu, A., and Nepal, S. DeepTheft: Stealing DNN model architectures through power side channel, September 2023. URL <https://arxiv.org/abs/2309.11894>. Published at 2024 IEEE Symposium on Security and Privacy (S&P).
- Garg, S., Goel, A., Jha, S., Mahloujifar, S., Mahmoody, M., Policharla, G.-V., and Wang, M. Experimenting with zero-knowledge proofs of training. In *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2023. doi: 10.1145/3576915.3623202.
- Heim, L. and Koessler, L. Training compute thresholds: Features and functions in AI regulation, 2024. URL <https://arxiv.org/abs/2405.10799>.
- Horowitz, M. 1.1 computing’s energy problem (and what we can do about it). In *2014 IEEE International Solid-State Circuits Conference (ISSCC) Digest of Technical Papers*, pp. 10–14, 2014. doi: 10.1109/ISSCC.2014.6757323.
- Hu, X., Liang, L., Li, S., Deng, L., Zuo, P., Ji, Y., Xie, X., Ding, Y., Liu, C., Sherwood, T., and Xie, Y. DeepSniffer: A DNN model extraction framework based on learning architectural hints. In *Proceedings of the 25th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pp. 385–399. ACM, 2020.
- Jia, H., Yaghini, M., Choquette-Choo, C. A., Dullerud, N., Thudi, A., Chandrasekaran, V., and Papernot, N. Proof-of-learning: Definitions and practice. In *2021 IEEE Symposium on Security and Privacy (SP)*, 2021. doi: 10.1109/SP40001.2021.00106.
- Latif, I., Newkirk, A. C., Carbone, M. R., Munir, A., Lin, Y., Koomey, J., Yu, X., and Dong, Z. Single-node power demand during AI training: Measurements on an 8-GPU NVIDIA H100 system. Technical report, Office of Scientific and Technical Information (OSTI), 2025. URL <https://www.osti.gov/biblio/2555926>.
- Maia, H. T., Xiao, C., Li, D., Grinspun, E., and Zheng, C. Can one hear the shape of a neural network? Snooping the GPU via magnetic side channel. In *31st USENIX Security Symposium (USENIX Security 22)*, pp. 4383–4400. USENIX Association, 2022.
- Monfared, S. K., Ganji, F., Holcomb, D., and Tajik, S. Timing and memory telemetry on GPUs for AI governance, 2026. URL <https://arxiv.org/abs/2602.09369>.
- Naghibijouybari, H., Neupane, A., Qian, Z., and Abu-Ghazaleh, N. Rendered insecure: GPU side channel attacks are practical. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, pp. 2139–2153. ACM, 2018.
- NVIDIA Corporation. NVIDIA A100 tensor core GPU architecture. Whitepaper, 2020. URL <https://www.nvidia.com/content/dam/en-zz/Solutions/Data-Center/nvidia-ampere-architecture-whitepaper.pdf>. Dense tensor-core peaks: 19.5 TFLOP/s fp32, 156 TFLOP/s tf32, 312 TFLOP/s fp16/bf16, 624 TOPS int8.
- O’Gara, A., Kulp, G., Hodgkins, W., Petrie, J., Immler, V., Aysu, A., Basu, K., Bhasin, S., Picek, S., and Srivastava, A. Hardware-enabled mechanisms for verifying responsible AI development. In *ICML 2025 Workshop on Technical AI Governance (TAIG)*, 2025. URL <https://arxiv.org/abs/2505.03742>.
- Pallipadi, V. and Starikovskiy, A. The ondemand governor: Past, present, and future. In *Proceedings of the Linux Symposium*, volume 2, pp. 223–238, 2006.
- Patel, P., Choukse, E., Zhang, C., Goiri, Í., Warriar, B., Mahalingam, N., and Bianchini, R. POLCA: Power oversubscription in LLM cloud providers, 2023. URL <https://arxiv.org/abs/2308.12908>.
- Rahman, R. and Tajdari, S. Detecting hidden ML training with zero-overhead telemetry, 2026. URL <https://arxiv.org/abs/2606.19262>.
- Reuel, A., Bucknall, B., Casper, S., Fist, T., Soder, L., Aarne, O., Hammond, L., Ibrahim, L., Chan, A., Wills, P., Anderljung, M., Garfinkel, B., Heim, L., Trask, A., Mukobi, G., Schaeffer, R., Baker, M., Hooker, S., Solaiman, I., Luccioni, A. S., Rajkumar, N., Moës, N., Ladish, J., Bau, D., Bricman, P., Guha, N., Newman, J., Bengio, Y., South, T., Pentland, A., Koyejo, S., Kochenderfer, M. J., and Trager, R. Open problems in technical AI governance, 2025. URL <https://arxiv.org/abs/2407.14981>.
- Sastry, G., Heim, L., Belfield, H., Anderljung, M., Brundage, M., Hazell, J., O’Keefe, C., Hadfield, G. K., Ngo, R., Pilz, K., Gor, G., Bluemke, E., Shoker, S., Egan, J., Trager, R. F., Avin, S., Weller, A., Bengio, Y., and Coyle, D. Computing power and the governance of artificial intelligence, February 2024. URL <https://arxiv.org/abs/2402.08797>.
- Sevilla, J., Heim, L., Ho, A., Besiroglu, T., Hobbhahn, M., and Villalobos, P. Compute trends across three eras of machine learning. In *2022 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8, 2022. doi: 10.1109/IJCNN55064.2022.9891914.

- Shavit, Y. What does it take to catch a Chinchilla? Verifying rules on large-scale neural network training via compute monitoring, 2023. URL <https://arxiv.org/abs/2303.11341>.
- Tang, B. J., Chen, Q., Weiss, M. L., Frey, N., McDonald, J., et al. The MIT Supercloud workload classification challenge, 2022. URL <https://arxiv.org/abs/2204.05839>.
- The White House. Executive order 14110: Safe, secure, and trustworthy development and use of artificial intelligence. Federal Register 88 FR 75191, 2023. URL <https://www.federalregister.gov/documents/2023/11/01/2023-24283/>.
- The White House. Executive order 14148: Initial rescissions of harmful executive orders and actions. Federal Register Doc. 2025-01901, 2025. URL <https://www.federalregister.gov/documents/2025/01/28/2025-01901/>.
- Wei, J., Zhang, Y., Zhou, Z., Li, Z., and Al Faruque, M. A. Leaky DNN: Stealing deep-learning model secret with GPU context-switching side-channel. In *2020 50th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, pp. 125–137. IEEE, 2020.
- Xu, H., Gong, C., Liu, B., Zheng, H., Chen, B., and Li, M. Wave: Leveraging architecture observation for privacy-preserving model oversight. In *Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems (AS-PLOS)*, 2026. doi: 10.1145/3779212.3790247.
- Yang, Z., Adámek, K., and Armour, W. Part-time power measurements: nvidia-smi’s lack of attention, 2023. URL <https://arxiv.org/abs/2312.02741>.

A. From the CMOS power law to the exchange rate

Section 2.1 sketches the route from the CMOS power law to the exchange rate; the factor decomposition of Section 4.1 additionally adopts the factored model Equation (25). We give the derivation in full here and locate the terms the factorisation drops. Resolving the static term of Equation (1), total chip power separates into a dynamic (switching) part and a static part,

$$P = P_{\text{dyn}} + P_{\text{stat}}, \quad P_{\text{dyn}} = \alpha C V^2 f, \quad P_{\text{stat}} = V I_{\text{leak}} + P_{\text{fixed}}, \quad (32)$$

with α the activity factor, C the total capacitance of the logic (so αC is the capacitance switched per cycle), V the supply voltage, f the clock, $V I_{\text{leak}}$ the leakage power, and P_{fixed} the cooling, I/O, and idle draw. The forward model Equation (2), $P = rF + P_0$, is slope–intercept in the operation rate F , so the exchange rate is the *marginal* energy per operation and the overhead is the intercept,

$$r = \frac{\partial P}{\partial F}, \quad P_0 = P|_{F=0}. \quad (33)$$

Experiment 2 estimates Equation (33) directly, regressing mean power on sustained nominal-op rate with idle cells as the $F = 0$ anchors (Section 4.2).

The dynamic term gives r . The operation rate is $F = N_{\text{op}} f$, where N_{op} is the number of nominal operations the datapath retires per cycle at the operating format (an idle unit lowers the activity α , not the clock f). At a fixed format the clock is thus the knob through which the operation rate moves, $\partial f / \partial F = 1/N_{\text{op}}$, and the marginal rate of Equation (33) follows by the chain rule, holding α , C , and V fixed,

$$r = \frac{\partial P_{\text{dyn}}}{\partial F} = \frac{\partial P_{\text{dyn}}}{\partial f} \frac{\partial f}{\partial F} = \frac{\alpha C V^2}{N_{\text{op}}} = \alpha \left(\frac{C}{N_{\text{op}}} \right) V^2, \quad (34)$$

and the clock drops out: energy per operation does not depend on how fast the operations run, only on what each switching event costs and how many operations it delivers. The rate is constant in F , so $P_{\text{dyn}} = rF$ is linear through the origin, exactly the affine form the forward model Equation (2) assumes and Experiment 2 fits. The capacitance switched per operation, C/N_{op} , has two structural sources: the arithmetic datapath, whose width is set by the numeric format, and the data-movement path, set by where the operands reside and which kernel runs. Factoring it accordingly,

$$\frac{C}{N_{\text{op}}} = \kappa(\text{precision}) C_{\text{eff}}(\text{locality, kernel}), \quad (35)$$

with κ collecting the format scale (datapath width and the operation packing N_{op}) and C_{eff} the locality and kernel structure, and substituting into Equation (34), gives the model Equation (25). This is a re-parameterisation by physical origin, not a unique factorisation: κ , α , and C_{eff} are all switching quantities, grouped by cause—format, operand data, memory locality—so that each maps to a single verifier category.

Short-circuit sits in r , inside the switched capacitance. Chandrakasan et al. (1992) also carry a *short-circuit* term for the direct-path current that flows while an input transitions and the NMOS and PMOS networks briefly conduct together. It is triggered by the same switching events as P_{dyn} , so it scales with the operation rate and survives the marginal derivative of Equation (33). In a well-designed circuit it is a small fraction of the switching power, so we fold it into the effective switched capacitance C_{eff} of Equation (35): it lands in r , but introduces no new verifier category.

Leakage sits in P_0 , not r . The static power P_{stat} of Equation (32) flows whether or not the machine computes, so $\partial P_{\text{stat}} / \partial F = 0$: by Equation (33) it contributes nothing to the marginal rate r and is entirely the intercept P_0 , where the derivation of Equation (10) already places it and carries it as a bounded nuisance band. A leakage-per-operation term $r_{\text{leak}} = P_{\text{stat}}/F$ would appear only if r were defined as the *average* energy per operation, P/F , which mixes in the intercept and is operating-point dependent; we keep the marginal definition Equation (33), consistent with Equation (2), and so carry no separate leakage term in r .

B. What the bound requires of the measurements

Appendix A derives the affine forward map Equation (2) from device physics. The security bound does not need it. This appendix states the weaker condition it does need, and why the two band edges are estimated by different means.

Two roles, two estimands. The declared edges enter Equation (17) only through their *difference* $r_{hi}^{dec} - r_{lo}^{dec}$ at a common saturated rate, so any baseline amortised equally into both edges cancels and a measured total quotient E/C is the correct estimand for the numerator (the overhead enters once, separately, as ΔE_0). The covert cost enters as the *denominator* r_{lo}^{cov} , where nothing cancels: the bound is a valid security guarantee only if r_{lo}^{cov} is a genuine one-sided *lower* bound on the *incremental* covert energy. Write the true draw as $P = g_x(F_{dec}, F_{cov})$, some device-dependent function of what the silicon executes, with the subscript x standing for the execution details the verifier never sees. The condition is then

$$g_x(F_{dec}, F_{cov}) - g_x(F_{dec}, 0) \geq r_{lo}^{cov} F_{cov}, \quad (36)$$

for every admissible execution g_x . Under only this incremental condition the honest energy slack of Equation (17) still divides by r_{lo}^{cov} , so the floor holds without the global affine map of Equation (2); the affine model is the special case $g_x = P_0 + r_{lo}^{dec} F_{dec} + r_{lo}^{cov} F_{cov}$, which additionally makes $\beta_{en}(h)$ linear in h . We therefore identify r_{lo}^{cov} by a dose–response intervention (Experiment 3 for the generic int8 edge, Experiment 4 for the useful-operand edge), not by fitting a power–rate curve; a total-quotient covert edge is inadmissible here, and every covert edge the bound uses is a marginal dose slope.

Why the between-shape slope is not that estimand. Experiment 2’s regression of power on rate across kernel shapes is a consistency diagnostic for the affine map, never an input to the bound, and the two roles above say why. Those shapes move C_{eff} and α along with F , so the resulting spread in r is the same physical variation Experiment 1 maps deliberately as the exchange-rate band—swept on purpose there, drifting as a side effect here. A slope fitted through them therefore averages over shapes whose non-covert factors vary together, where Equation (36) needs the derivative at *fixed* declared work. That is why the covert edge is identified by a dose–response intervention at a fixed declared burst rather than read off a power–rate curve, and why the fit loosening across kernel classes ($R^2 = 0.99$ within the compute-bound cubes, 0.852 across all ten shapes) is expected rather than a failure of the model.

Three kinds of quantity, three kinds of status. The inputs to the bound do not share one statistical status, and we state each rather than claim a single coverage guarantee over all of them. The covert denominator—the one quantity a lower bound is genuinely sensitive to—is a one-sided 95% lower confidence limit on a regression slope (Experiments 3 and 4, and the useful-band dose slope), coverage-controlled by construction. The declared band $[r_{lo}^{dec}, r_{hi}^{dec}]$ enters only as a *difference* of per-cell-median edges at a common saturated rate, so the amortised baseline cancels and it acts as a band *width* rather than a point estimate needing its own coverage; its target population is the achievable per-shape operating points of the declared format on the audited A100 SKU, and the edges are that observed operating envelope. Used as design-time security edges they remain sampled operating points, not certified physical extrema. The overhead width ΔP_0 and the capacity and γ inputs are conservative worst-case *design* parameters, not sampled estimates at all: ΔP_0 is the observed idle envelope stretched by the adopted $\pm 5\%$ meter tolerance (widening it only widens β), and C_{max} and γ are published hardware peaks. Nothing here claims simultaneous coverage over the rectangle these edges span; that rectangle is a deliberately conservative over-approximation, since no single schedule attains its combined extrema (Assumption 3.1).

C. Two conservative accounting choices, quantified

Two accounting choices in the model bound are deliberately adversary-favorable; this appendix quantifies what each concedes.

Published peaks vs. achieved rates. The capacity clip Equation (20) is evaluated at the *published peak* capacities rather than the rates our kernels achieve. Peaks give the covert stream more room than it can really use, so the conceded β is larger than an achieved-rate evaluation would give. On our own measured peaks (246.9 Top/s fp16 and 307.7 Top/s int8, so $\gamma = 1.25$ rather than 2) the bare-meter rung falls from 1.95 to 1.22 and the precision rung from 1.16 to 0.91. Assumption 3.1 asks for achieved throughputs when it certifies that a particular *witness* is realisable in one window, and the witnesses of Section 5.1 are checked that way (Appendix D); the model bound itself is deliberately evaluated at peaks.

Baseline power charged twice on the bare-meter rung. Every entry of the bare-meter triple carries the estimand its algebraic role requires: the declared pair $(r_{hi}^{dec}, r_{lo}^{dec}) = (19.5, 0.786)$ pJ/op as total quotients entering through their difference, the covert entry a one-sided marginal lower bound (0.515 pJ/op). This is why r_{lo}^{dec} and r_{lo}^{cov} stay distinct even here, where they describe the same physical corner: they are a third apart precisely because a total quotient amortises baseline power—leakage, cooling, idle draw—that a marginal cost does not, so a single number could satisfy at most

one of the two roles, and getting the denominator’s role wrong understates β (Appendix B). Keeping them distinct costs little in the numerator, where the ceiling dominates the difference, but it matters in the denominator: the overhead term $\Delta P_0 / (r_{\text{lo}}^{\text{cov}} F_{\text{max}}) \approx 0.26$ is half again the 0.168 a total-quotient covert edge would give.

One conservatism does survive the bookkeeping. Baseline power is charged twice: once amortised inside the total-quotient ceiling $r_{\text{hi}}^{\text{dec}}$, and again as the overhead band ΔE_0 in its own term. Both charges widen the rung, so the figure we quote is an outer sensitivity rather than an exact width—the safe direction for a worst case, and small next to the rung’s dominant term.

D. The matched-energy witness protocol

The matched-pair witnesses of Section 5.1 hold the trusted declaration and the observation window fixed by construction: run A executes N declared fp16 iterations on dear operands and idles to a fixed $T = 10$ s window; run B executes the same N iterations on cheap operands plus a covert int8 fill tuned to match A’s energy, padded to the same window. This appendix records the protocol details and checks behind the accepted loads quoted in the body (0.19, 0.29, and 0.41 machines at $h = 0.10, 0.15$, and 0.216).

The energy match, and its margin. Over eight interleaved repeats at $h = 0.10$, run A drew 1291.3 ± 4.4 J and run B 1279.3 ± 5.6 J, matching to $|\Delta E|/E = 0.9\%$ while B hid 3.5×10^4 covert int8 iterations behind the one-sided 3σ no-alarm rule. Two features of the setup are conservative and one is thin. Clock locking is denied on the platform (Section 4.2), so the arms are matched by cadence rather than by a held clock: both reach the same 1410 MHz compute clock, the per-window *median* differing only through the idle-pad share, as in Experiment 3. Run A’s dear burst draws ≈ 500 W against run B’s ≈ 245 W and so sits at the card’s software power cap, which truncates A’s measured honest energy; since the threshold is set from A’s mean and the covert dose is tuned against it, the cap tightens the bar B must clear. The thin part is the topmost witness: at $h = 0.216$ the cheating arm’s spread is $2.6\times$ the honest arm’s, and its dearest repeat (1969.5 J) lands above the honest arm’s dearest (1964.6 J), clearing the threshold by 0.86% of the honest mean. The 0.41 figure is the largest load we could get accepted, not a comfortable one.

Achieved rates: lower witnesses, not corner attainment. The achieved incremental rates are 1.84 pJ/op for the dear declared arm (below the band ceiling 2.19), 0.82 for the cheap declared arm, and 0.54 for the covert fill (above the marginal edge 0.51): the witnesses meet the shared constraints of Assumption 3.1 but do not jointly reproduce the box’s extreme edges, which is why the accepted load 0.41 at $h = 0.216$ sits below the model’s $\beta(0.216) = 0.72$.

The busy-time ceiling. Run B multiplexes cheap declared work and roughly two covert operations per declared operation into one window, so above $h_{\text{max}} \approx 0.24$ its work no longer fits T . The ceiling is a property of this construction, not a general feasibility bound: it is set by this kernel’s achieved rates and covert dose ratio, and includes a 10% protocol margin so the idle pad never underflows (without it, $h_{\text{max}} = 0.27$ and the restricted worst case of Section 5.1 is 0.82 rather than 0.77). A different covert workload would move it.

E. Benchmark protocol, hardware, and provenance

Experiment 1 (Section 4.2) reads the *total energy quotient* E/C off a single A100 by holding the workload fixed and moving one factor of Equation (25) at a time. The quotient is the band-edge estimand of Section 4.2—not the marginal rate $r = \partial P / \partial F$ of Equation (2), which Experiments 3 and 4 estimate as regression slopes—and we keep the two apart throughout. Both declared band edges are read at the same saturated rate, so their amortised baselines cancel in the band-edge difference the floor uses; the $\approx 5\%$ spread in achieved rate across cells leaves a small residual, which widens the slack conservatively. This appendix gives the hardware, provenance, and mechanics: the workload and why its operation count is exact, how each driver is set, and how the energy is metered. Experiments 2 and 3 reuse the same protocol with a different swept axis; Experiment 3 reuses its metering inside fixed-length composite windows (declared burst + covert dose + idle pad).

Hardware and provenance. Experiments 1–3 and 5 ran sequentially on one NVIDIA A100-SXM4-80GB (driver 595.84), in a single Slurm allocation on the OzSTAR facility on 22 July 2026 (340, 85, 70, and 135 recorded windows respectively). The released records and capture scripts index experiments by capture identifier, which does *not* match the numbering used here: the captures are named in the order they were written, the experiments in the order the argument needs them. The

mapping is Experiment 1 $\rightarrow$ `exp1_exchange`, Experiment 2 $\rightarrow$ `exp3_overhead`, Experiment 3 $\rightarrow$ `exp7_marginal`, Experiment 4 $\rightarrow$ `qblock_marginal`, and Experiment 5 $\rightarrow$ `exp2_operand`. Two identifiers carry no experiment number at all: `exp4_core` is the cross-device sweep and `qblock` the trained-weight total quotient reported only as a contrast to Experiment 4’s marginal edge. Read the identifiers, not the digits in them; the repository’s README carries the same table. Every record carries the device UUID, sweep id, configuration snapshot, and git commit, and the band-extraction code refuses to mix devices silently; because Experiment 3’s marginal cost and Experiment 1’s declared band price the same floor, the guard requires them to share one device. Seven cards of one SKU appear in total, by UUID prefix: GPU-ac1cbae9 (Experiments 1–3 and 5); GPU-4bb93dfb, GPU-617f2717, GPU-aa412071 (cross-device sweep); GPU-757ff02a (the trained-weight total quotient, reported only as a contrast to the marginal edge) and GPU-747045fa (Experiment 4, the incremental useful-work edge that prices the bottom two rungs of Section 5); GPU-108c1549 (the matched-energy witnesses of Section 4.2). Each single-device band is therefore compared against declared edges measured on a different card, and the same-SKU spread (2–3% on every band edge; Section 4.2) bounds the device variation it inherits. Full manifests accompany the released records; the trained weight/activation tensors themselves are not committed, but the manifests record their SHA-256 hashes and the GPU pre-flight validation metrics, so the extraction is checkable against the release but not replayable from the repository alone.

Operating point. The platform denied clock locking and power capping, so DVFS is a recorded axis rather than a controlled one: each window logs the governor-selected SM clock (achieved 1335–1410 MHz on the band sweeps), the bands are stratified into 45 MHz clock strata, and no window was flagged with a *harmful* throttle (thermal or hardware power brake). The densest operand bursts do sit at the card’s software power cap by design, which truncates their measured power and is therefore conservative for a dearest-rate edge.

Meter accuracy and cadence. `nvmlDeviceGetTotalEnergyConsumption` exposes a cumulative millijoule register integrated in firmware. NVML publishes no accuracy specification for this counter on GA100; the documented $\pm 5\%$ applies to the power readings of the earlier Fermi and Kepler generations, and we adopt that figure as a working meter tolerance wherever one is needed, always in the direction that widens the bound. The counter’s underlying power telemetry updates on a ~ 100 ms cadence and samples only part of the runtime (Yang et al., 2023); each measured cell therefore fills a ~ 1.5 s window ($\gg 100$ ms), as described under *Metering the energy* below.

The workload and its exact operation count. Every cell runs a dense square matrix multiply AB with $A \in \mathbb{R}^{M \times K}$, $B \in \mathbb{R}^{K \times N}$ (we take $M = N = K = n$). Each of the MN output entries is a dot product of length K — K multiplies and $K-1$ adds—so by the standard two-operations-per-multiply-add convention the matmul costs

$$C = 2MNK = 2n^3 \text{ op}, \quad (37)$$

counted as nominal operations of whatever format the cell runs (int8 multiply-adds included; Section 2.1), known *analytically* from the shapes alone, with no profiler or hardware counter involved. A measured window runs `iters` such matmuls back to back, so its nominal compute is $C_{\text{win}} = 2n^3 \text{ iters}$ and the energy per operation is $r = E_{\text{win}}/C_{\text{win}}$. Because C is exact under the nominal convention (the literal scalar count is $MN(2K-1)$, within 0.03% of $2MNK$ at the band cells’ $K = 2048$), r inherits the meter’s accuracy directly; and because r is intensive (energy *per* operation), cells of different size and `iters` compare directly.

Setting the swept drivers. The floor reads only the extremes of r , so what the sweep needs from each factor of Equation (25) is its *range*, not a typical value: every axis below is built to drive its factor to both extremes. Each must also move without disturbing the operation count Equation (37), or the quotient $r = E/C$ would change for two reasons at once—a constraint that dictates the design of the locality and operand axes in particular.

- **Precision** κ is the tensor dtype: `fp32`, `tf32` (fp32 operands routed through the tf32 tensor-core path), `fp16`, `bf16`, and `int8` (an integer-matmul kernel, `int8 × int8  $\rightarrow$  int32`). Lower precision charges less capacitance per multiply-accumulate, lowering κ .
- **Locality** C_{eff} is operand *residency* at a byte-identical shape, not problem size: sweeping it by changing n would change the operation count and confound the comparison, so instead we fix the shape at $n = 2048$ and vary how many distinct operand sets the window rotates through iteration to iteration. The A100 has a 40 MB L2; at $n = 2048$ an fp16 A – B operand set is ~ 16 MB, so:

- *resident* (one operand set, reused every iteration): the working set fits the L2, so most traffic is served on-chip and HBM activation is avoided: low C_{eff} ;
- *streamed* (eight operand sets, round-robin): a reuse is evicted before it recurs, so every iteration streams fresh tiles from HBM, engaging its arrays, peripheral logic, and I/O: high C_{eff} .

Both arms run the identical shape, kernel, and operation count Equation (37); only the memory-traffic path differs, so locality no longer confounds the count. The kernel stays compute-bound in both arms—a 2048³ matmul is high arithmetic intensity—so streaming raises C_{eff} without making the workload memory-bound, and Figure 5 accordingly shows locality moves r only slightly.

The third driver, supply voltage V^2 (DVFS), we could not sweep: the platform denied user clock-locking, so the clock was left to the governor and merely recorded per run. The operand *values* are a fourth axis: Experiment 1 crosses all six operand distributions (Appendix F) with precision and locality in a full factorial, so the pooled envelope already contains the operand spread; Experiment 5 re-runs the operand axis alone, at fixed precision and locality, as the descriptive operand-only map of Appendix F.

Metering the energy. Energy is read from the on-die NVML total-energy counter: the window energy is one subtraction $E_1 - E_0$, with a `cuda.synchronize()` fence bracketing each latch so the delta spans exactly the window’s kernels; where the counter is unavailable we integrate the power query in a background sampler. Each cell fills a ~ 1.5 s window with matmuls run back to back on pre-allocated buffers, with no host traffic inside the window: long enough that counter quantisation is negligible, and saturated enough that the joules are the datapath doing arithmetic, not allocation or transfer. A few warm-up windows absorb JIT and clock ramp, several measured windows follow, and the per-cell median is kept; shuffling the cell order and periodically re-measuring a fixed reference cell guards against slow thermal drift. The largest per-cell median over the sweep is $r_{\text{hi}}^{\text{dec}}$ with no format pinned, and the largest within one format is that format’s own ceiling; Figure 5 plots the full set.

F. How Experiment 5 varies the operands

Experiment 5 holds the kernel, shape, and precision fixed and moves *only the operand values*. Every configuration runs the identical matmul $A B$ with the identical operation count Equation (37), so the energy differences track the switching-activity factor α of Equation (25)—how many bits physically toggle in the datapath as the multiply–accumulates stream through. A datapath that sees the same bits over and over charges little capacitance; one whose bits flip on every operation charges the most. The six distributions are constructed to walk α from its floor to its ceiling.

Why this experiment is reported here and not in the body. It measures a width the body already has. Experiment 1’s factorial varies the same six operand families *inside every format*, so its fp16 cells contain this experiment’s comparison as one slice of a larger grid—and contain it at a matched clock, which this sweep does not have. Experiment 5 is therefore dominated: fp16-only, streamed-only, and confounded where Experiment 1 is balanced. We keep it for two reasons, neither of them load-bearing. It corroborates w_α from independent data, and it resolves the activity factor into named distributions, which is what the threat-model rungs of Section 5 act on when they exclude zero operands. No security bound rests on any number below.

The measurement. The same matmul at fixed kernel, shape, and precision (fp16), varying only the operand values. Operand values alone spread r by an observed factor $w_0 = 1.83$ (Figure 10), following the ordering zeros < low-entropy < declared-typical $\approx$ dense-random $\approx$ structured-sparse < adversarial-toggle. The clock confounds the comparison: locking is privilege-denied (Section 4.2) and the families did not run at a common clock—the cheapest (zeros, low-entropy) drew the highest clocks and the dearest the lowest, a monotone 1410 $\rightarrow$ 1324 MHz slide against operand cost. Experiment 1’s operand families, by contrast, share a median 1410 MHz.

That confound predicts its own direction, which is what makes the corroboration worth something. The cheap families ran fastest, so a common clock would push their energies up and spread the distribution *further*: w_0 should sit *below* a matched-clock measurement of the same factor. It does—1.83 against Experiment 1’s $w_\alpha = 1.88$, agreeing to 3% with the residual on the predicted side. Restricting Experiment 1 to the fp16 streamed cells that match this sweep’s configuration exactly gives 1.94, bracketing w_0 the same way. We offer the direction argument as a bias note rather than a proof: DVFS




Figure 10. The shape of the activity factor α (Experiment 5, $n = 96$ windows across six operand distributions). Kernel, shape, and precision (fp16) are held fixed and only the operand values vary; the clock was governor-selected, not locked, and slides monotonically against operand cost, so this map is descriptive. The shaded span is the observed spread $w_0 = 1.83$, just below Experiment 1’s matched-clock $w_\alpha = 1.88$, on the side the confound predicts. The minimum (zeros, dashed) is the descriptive fp16 operand floor 1.16 pJ/op as a total quotient—not the covert denominator, which is Experiment 3’s marginal edge.

voltage scaling, throughput, and baseline-power amortisation pull in competing directions, and one-sided uncertainty cannot repair an uncontrolled confound.

The minimum-energy run (zeros) puts the fp16 operand floor at 1.16 pJ/op as a total quotient—a descriptive edge only; the covert denominator the floor actually uses is Experiment 3’s marginal cost (a one-sided 95% lower confidence limit), which strips the baseline power this quotient amortises in.

The distributions. Below we show each on a 4×4 patch; the real operands are the same patterns at full size, generated by `operand_fill` in `powertoflops/bench/workload.py`.

1. Zeros—the activity floor. Every entry is zero, so no bits ever change and α is at its minimum:

$$A_{\text{zeros}} = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix}.$$

This is the cheapest run at fixed fp16 and sets the descriptive fp16 operand floor, 1.16 pJ/op as a total quotient; combined with the int8 format in Experiment 1’s factorial it reaches the cheapest measured quotient, 0.786 pJ/op, which is r_{lo}^{dec} with no format pinned. (A declared entry. The covert entry r_{lo}^{cov} is Experiment 3’s marginal cost, not either of these quotients; Section 4.2 and Table 1.)

2. Low-entropy—a few distinct levels. Entries are drawn from only four values $\{0, 1, 2, 3\}$. With so few distinct bit patterns, few bits differ between neighbouring operands, so operand-bit entropy—and hence switching—stays low:

$$A_{low-entropy} = \begin{pmatrix} 1 & 0 & 3 & 1 \\ 2 & 2 & 0 & 3 \\ 0 & 1 & 1 & 2 \\ 3 & 0 & 2 & 1 \end{pmatrix}.$$

3. Structured-sparse—half the operands forced to zero. Dense random values, but a fixed half of the columns are deterministically zeroed (every other column), giving exactly 50% zeros in a regular pattern: the structured sparsity real accelerators exploit, rather than random dropout. The persistent zero columns roughly halve the switching relative to a fully dense operand:

$$A_{sparse} = \begin{pmatrix} 0 & 0.4 & 0 & -0.9 \\ 0 & -0.7 & 0 & 0.2 \\ 0 & 0.8 & 0 & 0.5 \\ 0 & -0.3 & 0 & -0.6 \end{pmatrix}.$$

4. Dense-random—generic busy data. Every entry is an independent uniform draw from $[-1, 1]$: no zeros, full-width mantissas, high but not maximal switching. This is the realistic-workload reference:

$$A_{dense} = \begin{pmatrix} 0.31 & -0.82 & 0.55 & -0.14 \\ -0.67 & 0.09 & -0.48 & 0.73 \\ 0.21 & 0.64 & -0.95 & 0.38 \\ -0.52 & -0.27 & 0.86 & -0.41 \end{pmatrix}.$$

5. Declared-typical—the realistic-activation stand-in. Every entry is an independent standard-normal draw, $\mathcal{N}(0, 1)$: a proxy for the activations and weights a genuine declared workload streams, with full-width mantissas and no engineered structure. Its switching sits alongside dense-random—high but not maximal—and it is the operand class that characterises the honest-window energy scatter used by the declared-work model:

$$A_{decl} = \begin{pmatrix} 0.42 & -1.13 & 0.27 & -0.68 \\ -0.35 & 0.91 & -1.42 & 0.11 \\ 1.06 & -0.24 & 0.58 & -0.83 \\ -0.72 & 0.39 & -0.16 & 1.24 \end{pmatrix}.$$

6. Adversarial-toggle—the activity ceiling. Engineered for maximal switching: the matrix is a checkerboard of two *complementary* bit patterns, $0x5555\dots$ and its bitwise inverse $0xA555\dots$, so adjacent operands differ in *every* bit (maximal Hamming distance) and the operand path sees maximal input-level switching per operation. For floats the pattern is built in an integer bit-view of the same width and reinterpreted as the float type, so the raw bits—not the numeric value—are controlled exactly. Writing the two patterns as p and $\bar{p}$:

$$A_{adv} = \begin{pmatrix} p & \bar{p} & p & \bar{p} \\ \bar{p} & p & \bar{p} & p \\ p & \bar{p} & p & \bar{p} \\ \bar{p} & p & \bar{p} & p \end{pmatrix}, \quad p = 0x5555\dots, \quad \bar{p} = 0xA555\dots$$

A matmul forms each output as a dot product $out[i, j] = \sum_k A[i, k] B[k, j]$, streaming $k = 0, 1, 2, \dots$ and loading one operand into the multiplier’s input register each cycle. Because every row and every column of the checkerboard alternates $p, \bar{p}, p, \dots$, that register walks

$$\underbrace{0101\ 0101}_{k=0} \rightarrow \underbrace{1010\ 1010}_{k=1} \rightarrow \underbrace{0101\ 0101}_{k=2} \rightarrow \dots$$

flipping *all* its bits on *every* cycle. Dynamic CMOS energy is paid per bit-transition, so this maximises the switching the operand registers can drive: a w -bit operand toggles w bits per cycle, against $\approx w/2$ for dense-random (two random words differ in about half their bits) and 0 for zeros or any constant operand (the register never changes). The checkerboard is simply the 2-D layout that makes the row sequence (feeding A) and the column sequence (feeding B) both alternate this way at once, for every output entry.

The energy-per-operation spread from the cheapest patch (A_{zeros}) to the densest operands (dense-random, declared-typical, and the adversarial checkerboard), at fixed operation count, is the observed operand spread $w_0 = 1.83$ of Figure 10. The measured ordering follows the textbook one, zeros < low-entropy < structured-sparse < dense-random $\approx$ declared-typical $\approx$ adversarial-toggle. Per the clock confound of Section 4.2, the spread is descriptive within-device evidence, not a certified population edge—and maximal input-level toggling does not itself certify maximal switching inside the Tensor Core internals, whose operand encodings are opaque (nor do zero operands mean no hardware bits switch); no security bound in the paper rests on it.

G. Scratch space

Back of the envelope (superseded by Section 4.2). The order-of-magnitude estimate this paper carried before the A100 sweep existed. It is kept only as a record that the measured widths landed where device physics said they would; the two footnotes it used to carry have moved into the body, where the numbers they source now live.

Taking the factors in order: the precision scale κ spans roughly an order of magnitude from fp32 down to int8 (Horowitz, 2014) (our own sweep in Section 4.2 finds $13.7\times$ from fp32 to int8 at dense operands); the operand activity α adds a data-dependent factor of order two (no textbook number exists; we measure it directly in Section 4.2); locality and kernel set C_{eff} , for which we conservatively count no extra factor; and DVFS moves V^2 by a factor of about 1.5. Multiplied, the dearest per-operation cost the silicon reaches sits conservatively in the tens times the cheapest—but that span crosses numeric formats, and no single declared workload can be honestly explained at one format’s dearest cells while its capacity is priced at another’s (Section 3.2). So the widest declared band, the bare-meter setting, is a factor of tens; pin the format and the precision factor drops out, leaving operands and DVFS, for a declared band of a factor 2–3. That width is a numerator, and the slope s divides it by $r_{\text{lo}}^{\text{cov}}$, which the same physics puts a further factor below the declared floor; so s is of order one and the ambiguity term reaches, not dwarfs, one machine. The overhead term, which we cannot fix from first principles (Section 4.1), is additive and only raises the slack. On paper, then, energy binds at low utilisation, the capacity clip of Equation (20) takes over at high utilisation, and the worst-case floor sits at a substantial fraction of the machine—up to about its declared-format capacity as a model bound. A bare power meter certifies little about how many operations ran. We put numbers on this next by direct measurement.

Note that β is a floor, not a measurement error. It is the ambiguity that survives a perfect meter: because distinct workloads draw identical power, no reading of the trace can tell them apart because the data hold no information that separates them. Averaging, longer observation, and better estimators all fight measurement noise, and none of them lowers β .